

# Today: Outline

- **Neural networks:** artificial neuron, MLP, sigmoid units; neuroscience inspiration, output vs hidden layers; linear vs nonlinear networks;
- **Feed-forward networks**
- **Reminders:** First lab session tomorrow (Fri Sept 18)

# Fei-Fei Li



- Professor, Computer Science, Stanford University
- Co-Director of Stanford's Human-Centered AI Institute
- Previously Vice President at Google and Chief Scientist of AI/ML at Google Cloud
- Co-founder and chairperson of the national non-profit AI4ALL
- Online Deep Learning Course
- **"First, we teach them see, then they help us to see better."**

# Image Captioning



*A young boy holding a baseball bat*



*A man riding a horse next to a building*





# Introduction to Neural Networks

## Motivation

# Recall: Logistic Regression

$$0 \leq h_{\theta}(x) \leq 1$$

$$h_{\theta}(x) = g(\theta^T x) = \frac{1}{1 + e^{-\theta^T x}}$$

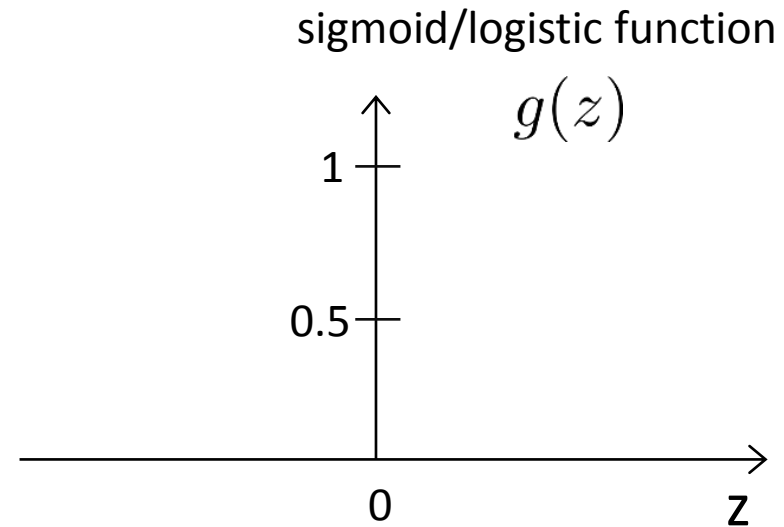
$$g(z) = \frac{1}{1 + e^{-z}}$$

Output is probability of label 1 given input

$$p(y = 1|x) = \frac{1}{1 + e^{-\theta^T x}}$$

predict “ $y = 1$ ” if  $h_{\theta}(x) \geq 0.5$

predict “ $y = 0$ ” if  $h_{\theta}(x) < 0.5$



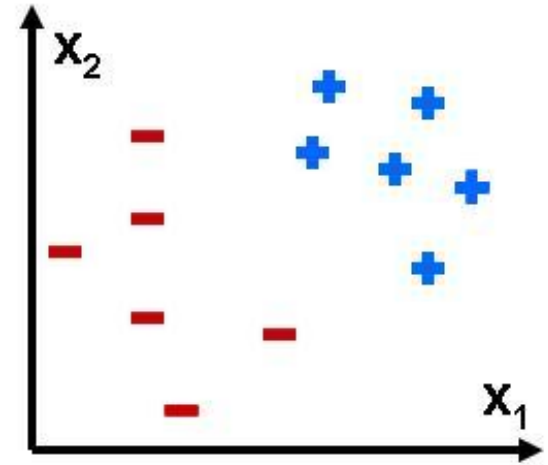
# Recall: Logistic Regression Cost

Logistic Regression Hypothesis:

$$h_{\theta}(x) = g(\theta^T x) = \frac{1}{1 + e^{-\theta^T x}}$$

$\theta$ : parameters

$D = \{x^i, y^i\}$ : data

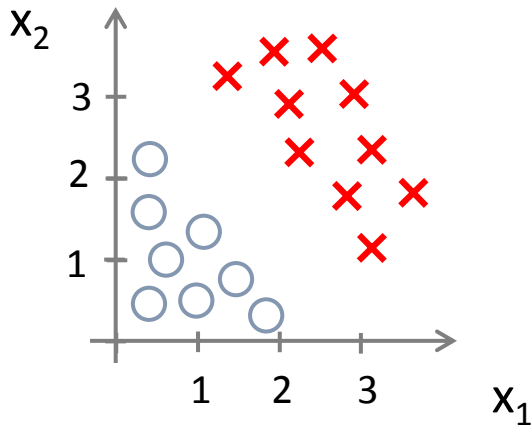


Logistic Regression Cost Function:

$$\begin{aligned} J(\theta) &= \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_{\theta}(x^{(i)}), y^{(i)}) \\ &= -\frac{1}{m} \left[ \sum_{i=1}^m y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log (1 - h_{\theta}(x^{(i)})) \right] \end{aligned}$$

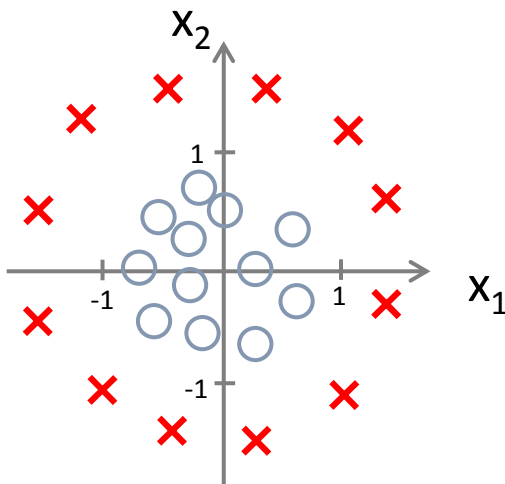
Goal: minimize cost  $\min_{\theta} J(\theta)$

# Decision boundary



$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2)$$

## Non-linear decision boundaries



Replace features with non-linear functions  
e.g. log, cosine, or **polynomial**

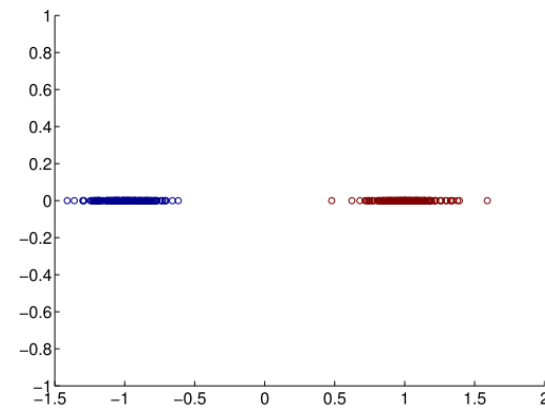
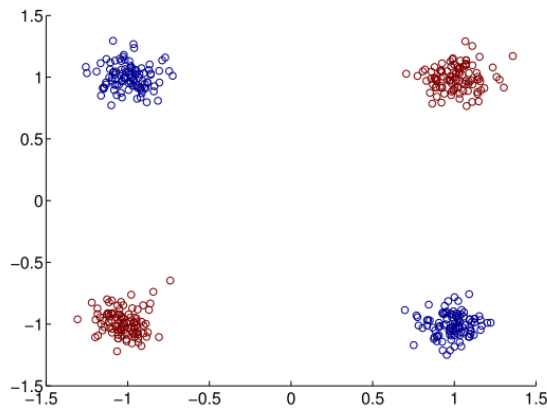
$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_2^2)$$

# Nonlinear basis functions

**Transform the input/feature**

$$\phi(x) : x \in R^2 \rightarrow z = x_1 \cdot x_2$$

**Transformed training data: linearly separable!**

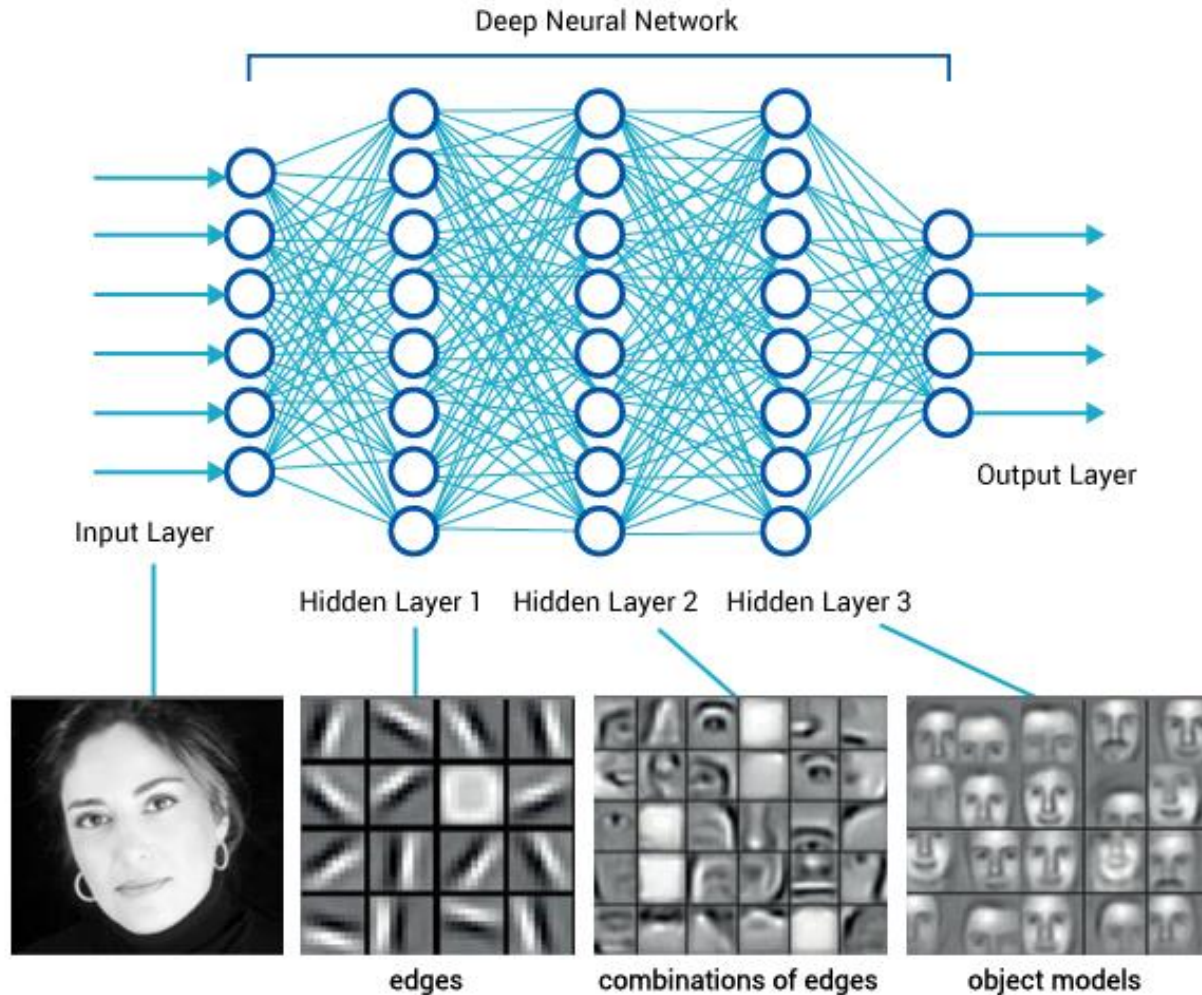




# Limitations of linear models

- Logistic regression and other linear models cannot handle nonlinear decision boundaries
  - Must use non-linear feature transformations
  - Up to designer to specify which one
- Can we instead **learn** the transformation?
  - Yes, this is what neural networks do!
- A **Neural network** chains together many layers of “neurons” such as logistic units (logistic regression functions)

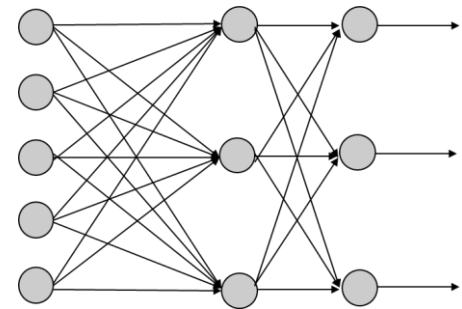
# Neural Networks learn features



# Neurons in the Brain

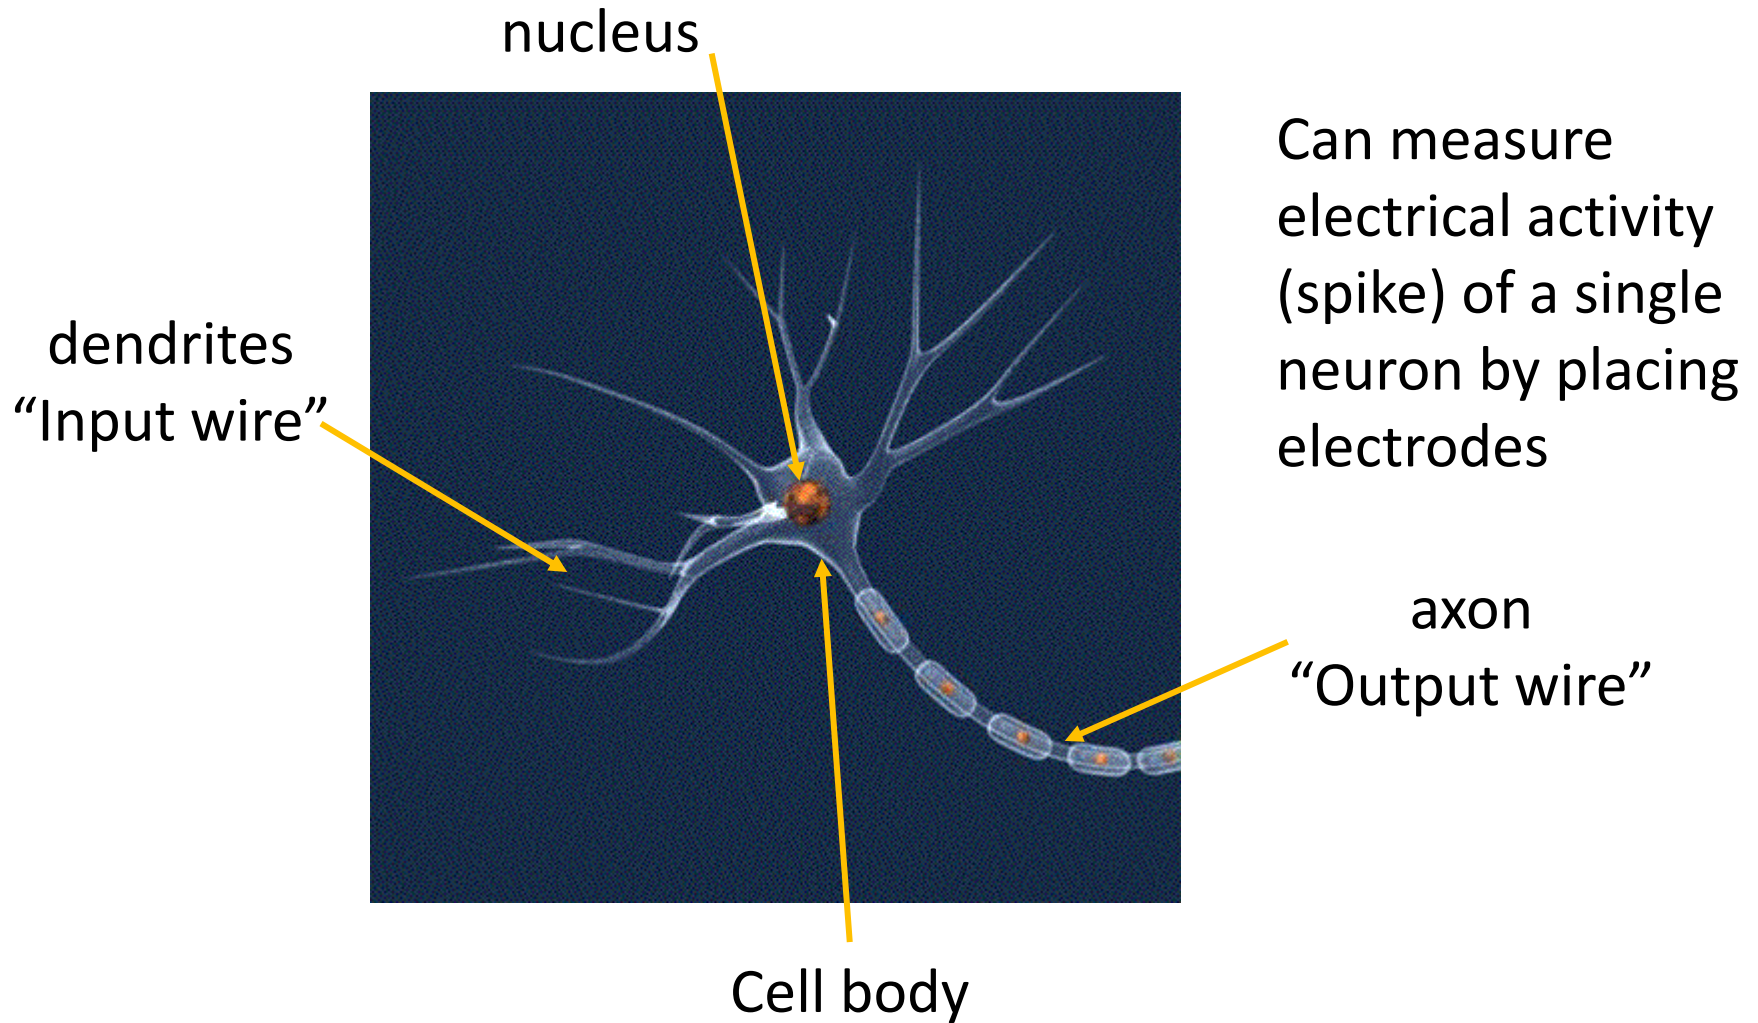


Inspired “Artificial Neural Networks”



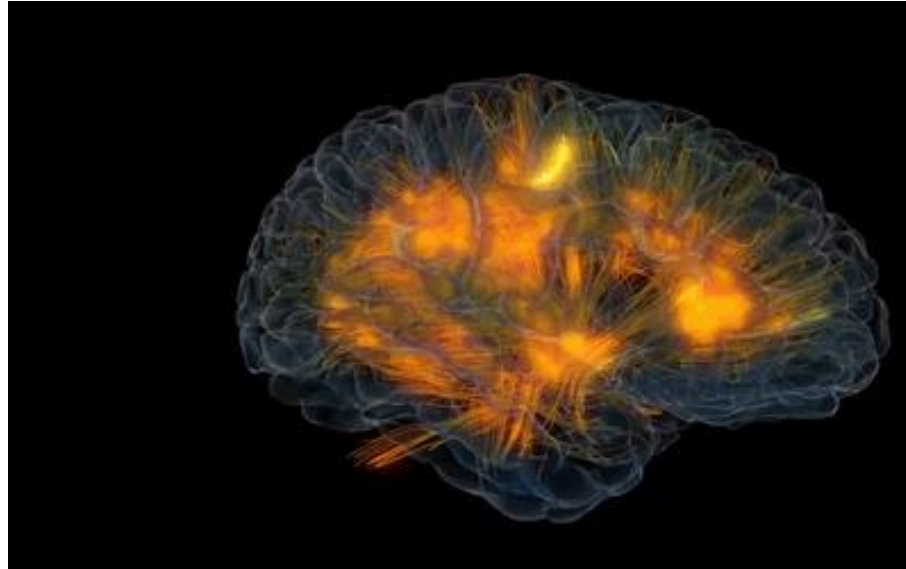
Neurons are cells that process chemical and electrical signals and transmit these signals to neurons and other types of cells

# Neuron in the brain



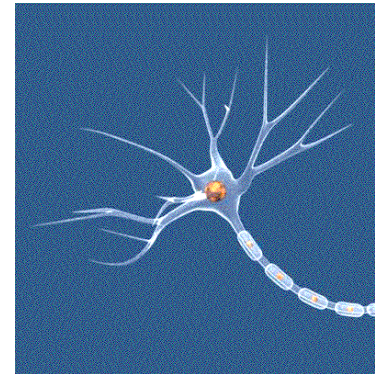
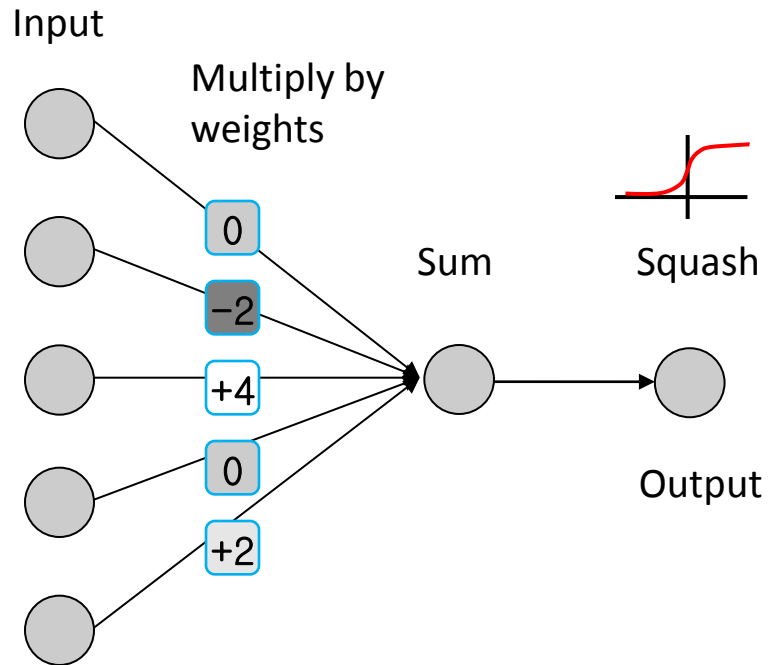


# Neural network in the brain



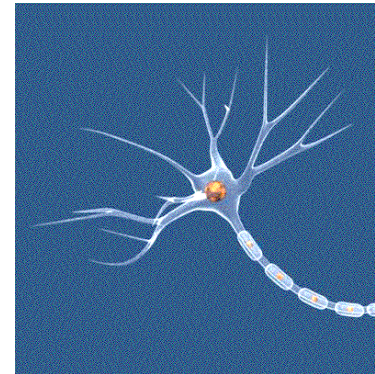
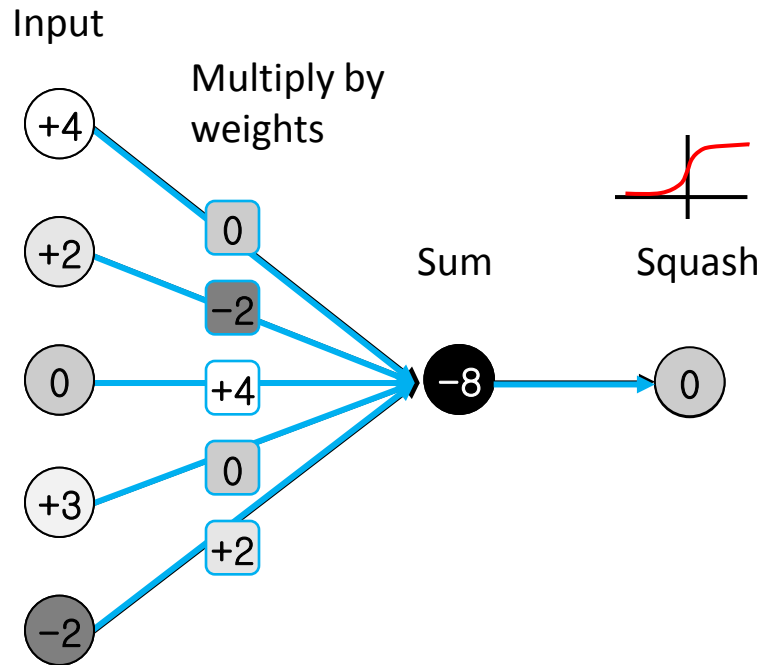
- **Micro networks:** several connected neurons perform sophisticated tasks: mediate reflexes, process sensory information, generate locomotion and mediate learning and memory.
- **Macro networks:** perform higher brain functions such as object recognition and cognition.

# Logistic Unit as Artificial Neuron



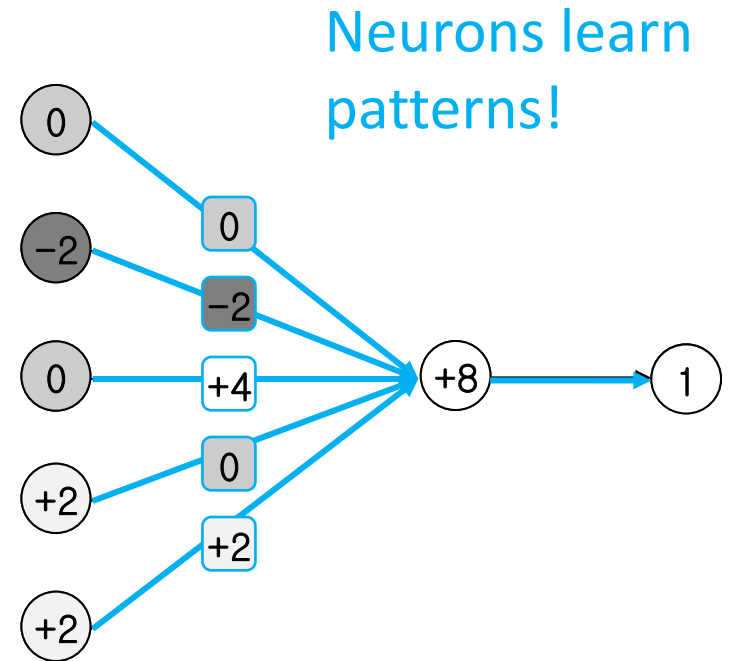
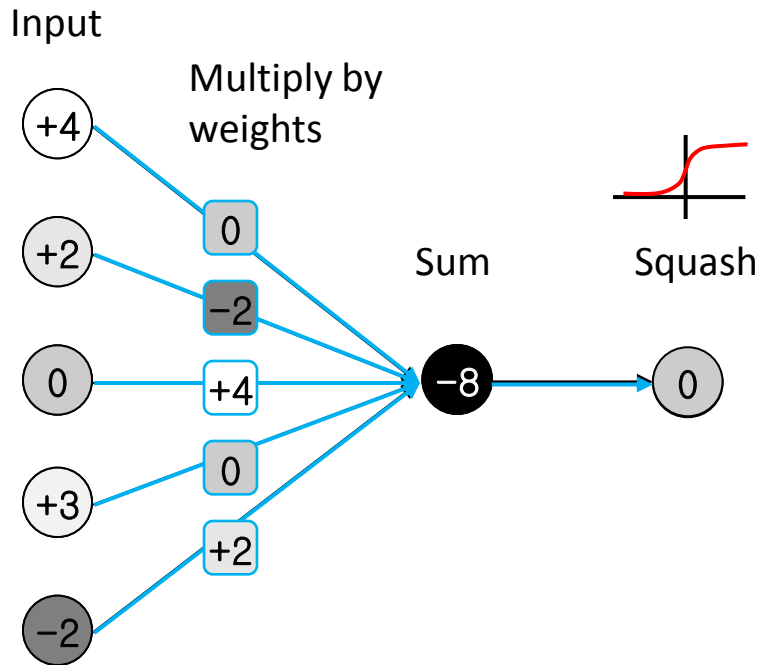


# Logistic Unit as Artificial Neuron



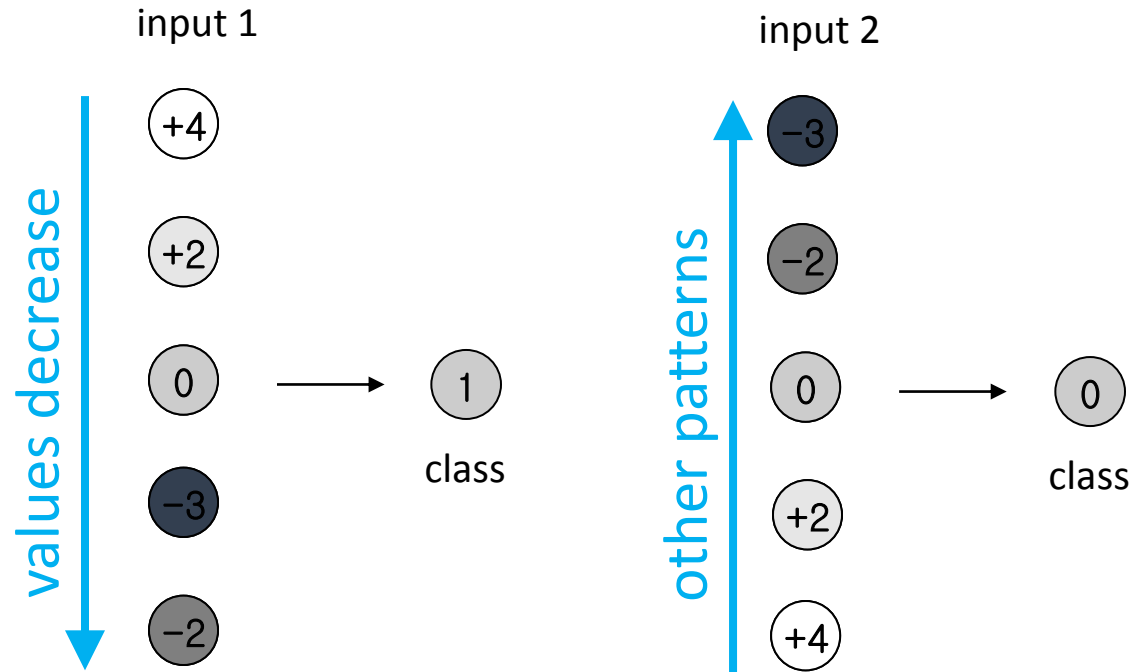
$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_n x_n$$

# Logistic Unit as Artificial Neuron



# Artificial Neuron Learns Patterns

- Classify input into class 0 or 1
- Teach neuron to predict correct class label
- Detect presence of a simple “feature”



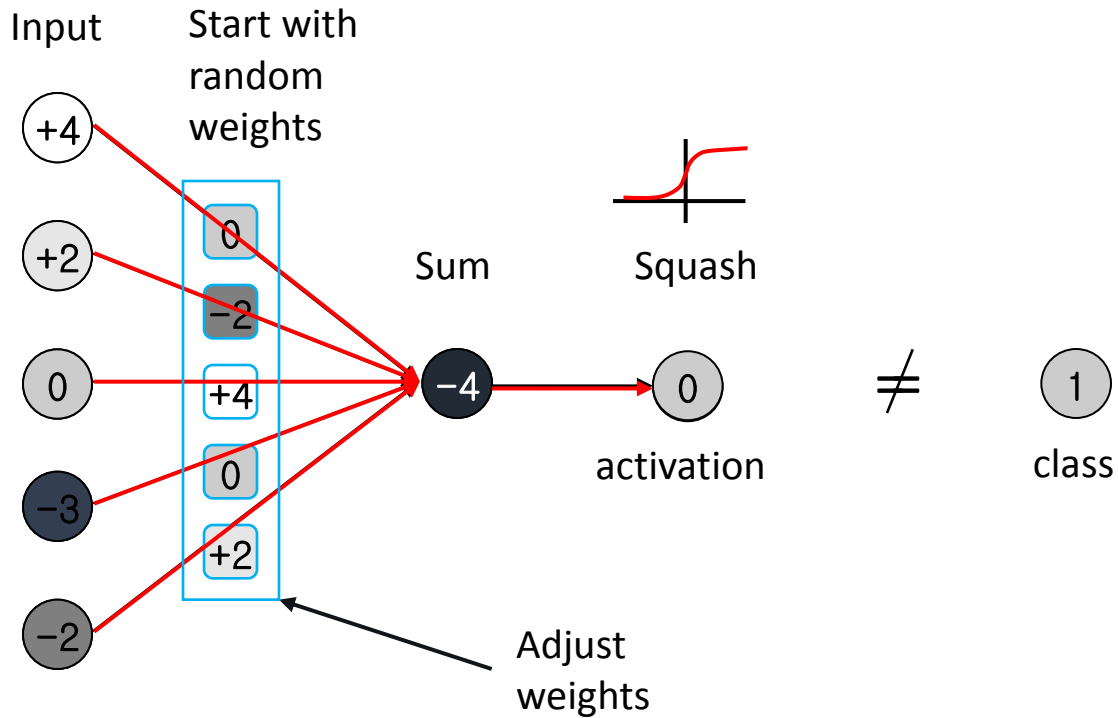
Example



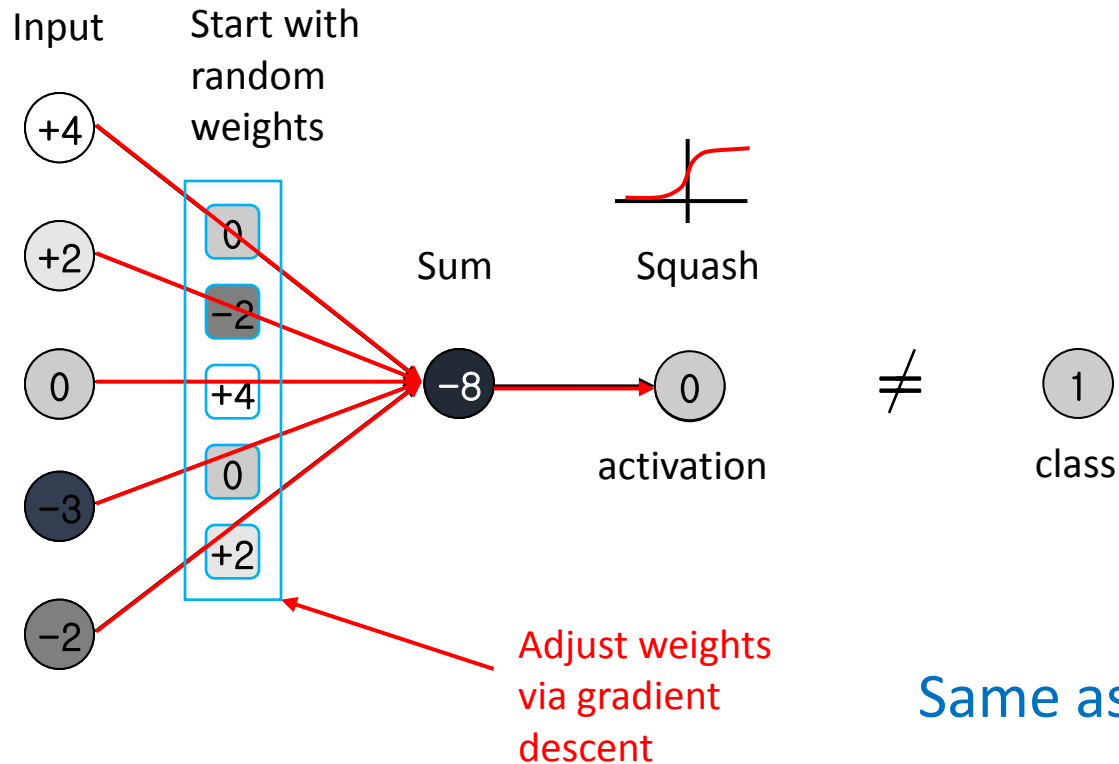
# Neural Networks: Learning

Intuition

# Artificial Neuron: Learning



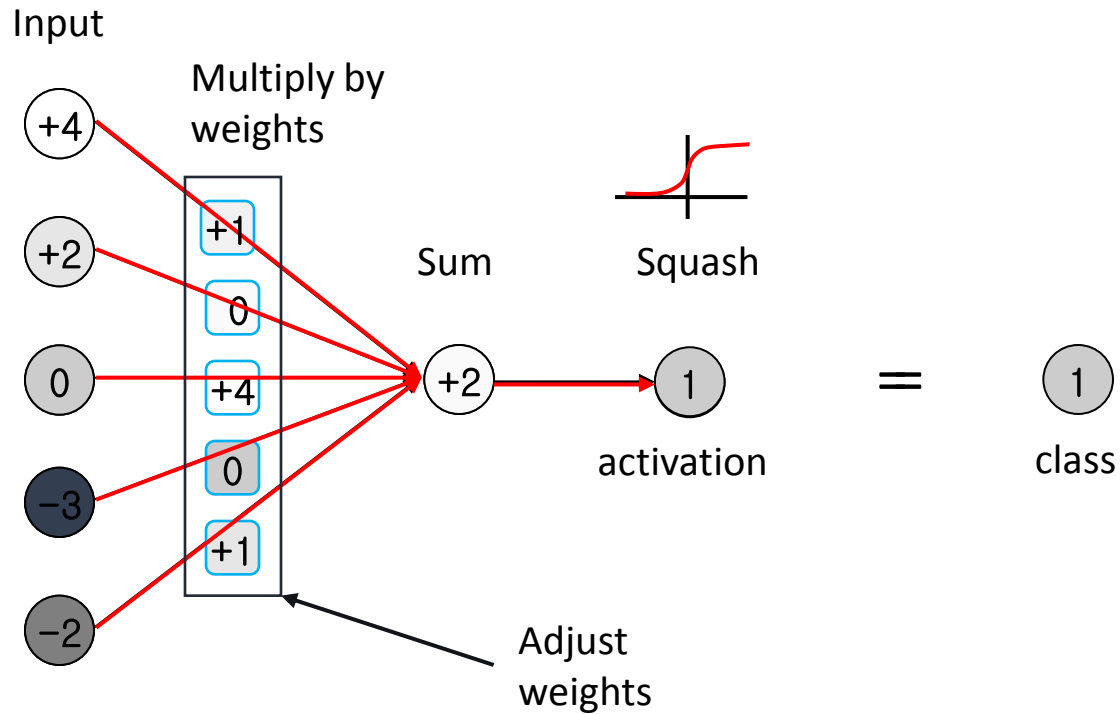
# Artificial Neuron: Learning



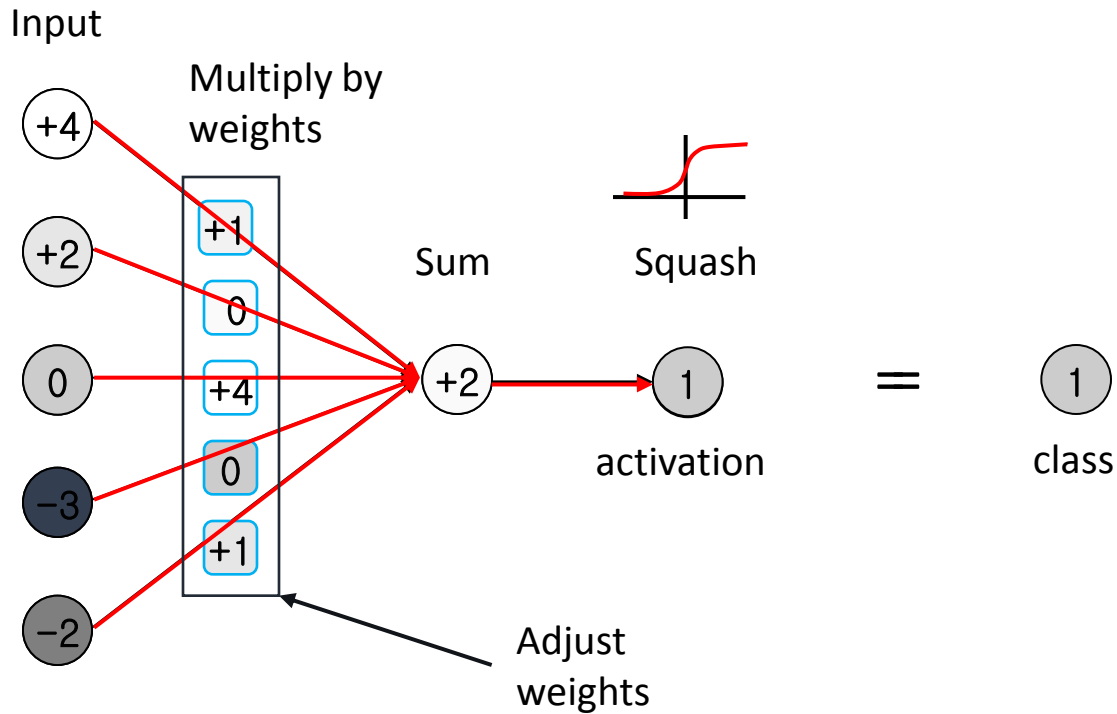
Same as in logistic regression



# Artificial Neuron: Learning



# Artificial Neuron: Learning



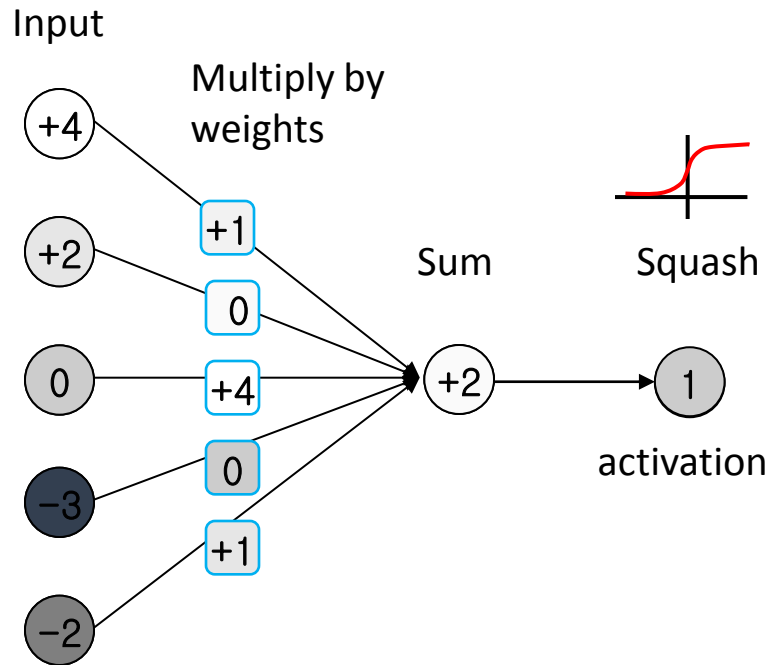
***Forward propagation of information through a neuron***



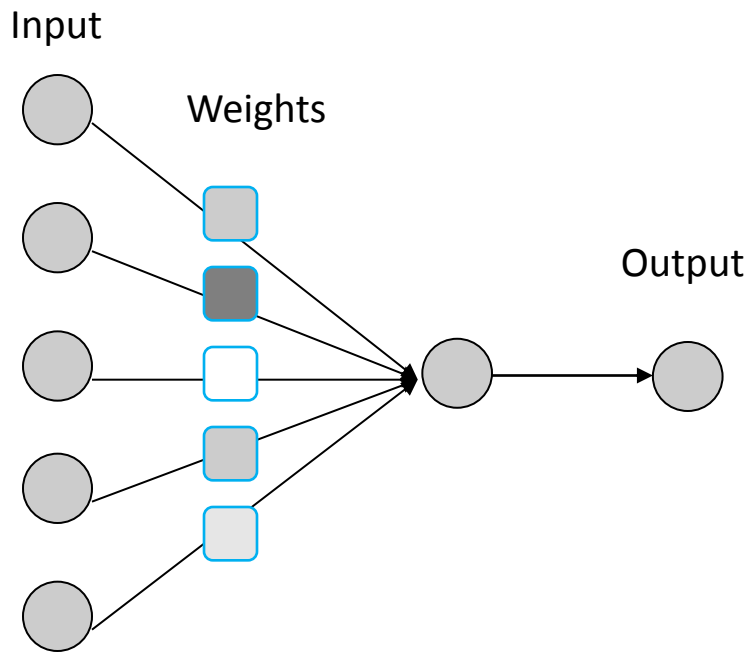
# Neural Networks: Learning

Multi-layer network

# Artificial Neuron: simplify

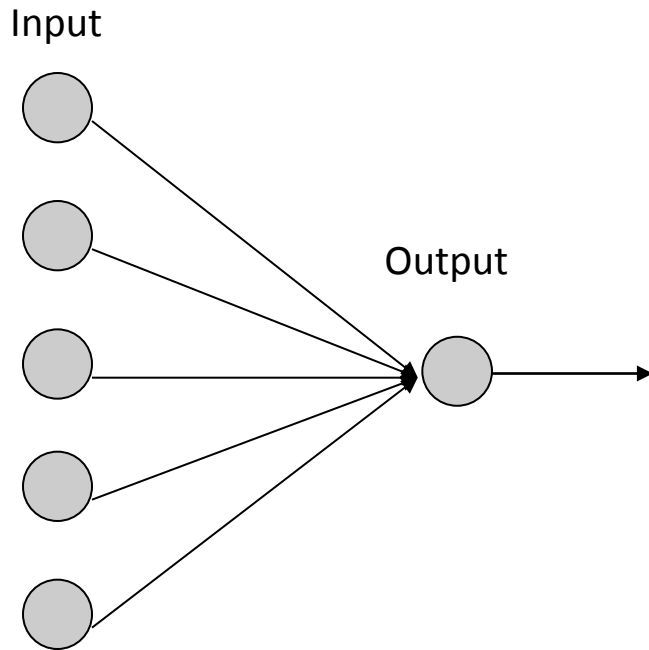


# Artificial Neuron: simplify



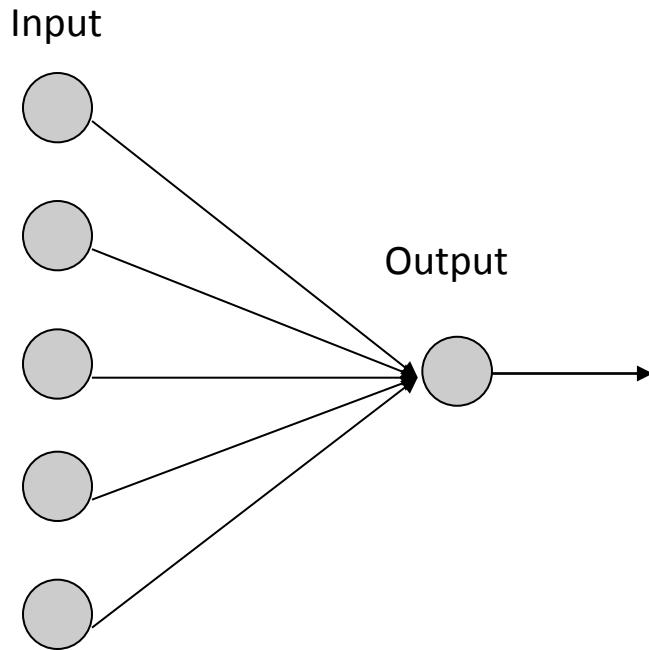
*A single neuron is also called a perceptron*

# Artificial Neuron: simplify

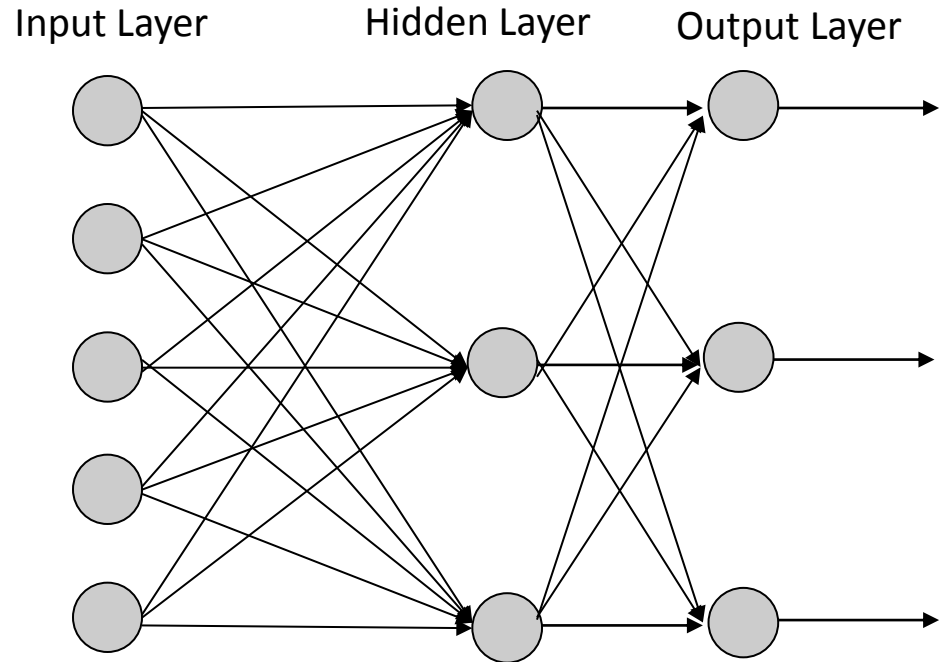




# Artificial Neural Network



Single Neuron

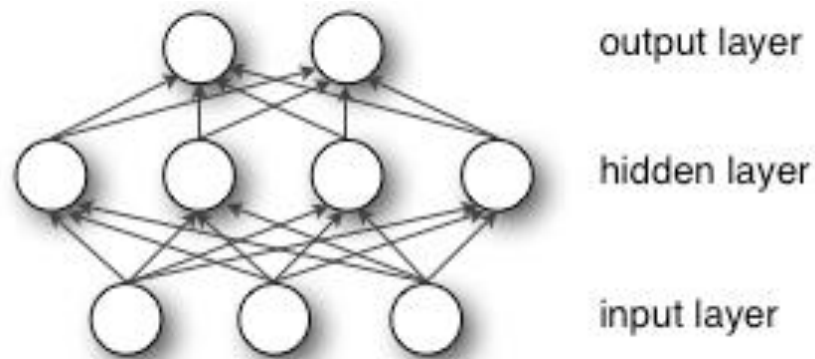


Neural Network (fully connected)

Deep Network: many hidden layers

# Multi-layer perceptron (MLP)

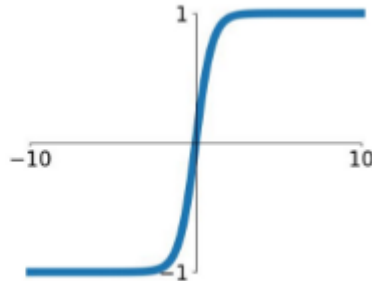
- Just another name for a feed-forward neural network
- Logistic regression is a special case of the MLP with no hidden layer and sigmoid output.



# Other Non-linearities

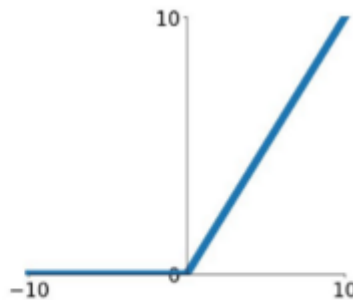
- Also called activation functions

**tanh**  
 $\tanh(x)$



$$\tanh(x) = \frac{2}{1+e^{-2x}} - 1$$

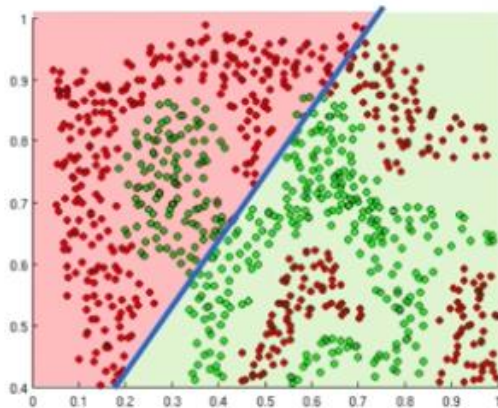
**ReLU**  
 $\max(0, x)$



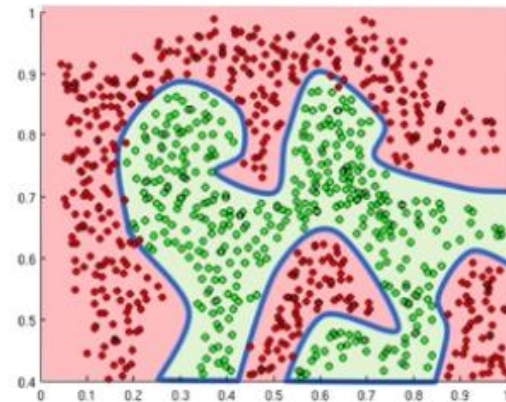
$$ReLU(x) = \begin{cases} 0 & \text{if } x < 0 \\ x & \text{if } x \geq 0 \end{cases}$$

# Importance of Non-linearities

*The purpose of activation functions is to **introduce non-linearities** into the network*



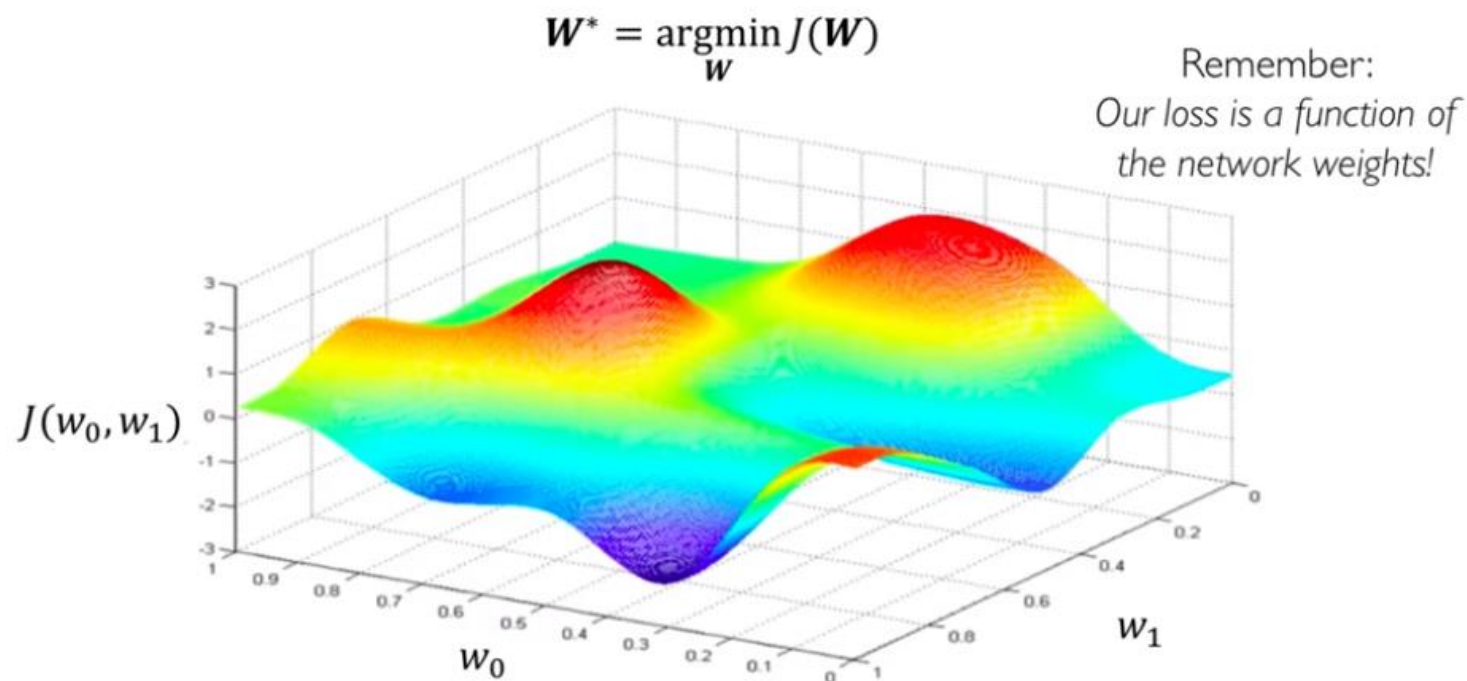
Linear activation functions produce linear decisions no matter the network size



Non-linearities allow us to approximate arbitrarily complex functions

# Loss Optimization

- Neural network parameters  $\theta$  are often referred to as weights  $W$ .



# Gradient Descent

## Algorithm

1. Initialize weights randomly  $\sim \mathcal{N}(0, \sigma^2)$
2. Loop until convergence:
3.     Compute gradient,  $\frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$
4.     Update weights,  $\mathbf{W} \leftarrow \mathbf{W} - \eta \frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$
5. Return weights




# Gradient Descent

## Algorithm

1. Initialize weights randomly  $\sim \mathcal{N}(0, \sigma^2)$
2. Loop until convergence:
3. Compute gradient,  $\frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$
4. Update weights,  $\mathbf{W} \leftarrow \mathbf{W} - \eta \frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$
5. Return weights

# Gradient Descent

## Algorithm

1. Initialize weights randomly  $\sim \mathcal{N}(0, \sigma^2)$
2. Loop until convergence:
3. Compute gradient,  $\frac{\partial J(W)}{\partial W}$  
4. Update weights,  $\mathbf{W} \leftarrow \mathbf{W} - \eta \frac{\partial J(W)}{\partial W}$
5. Return weights

# Gradient Descent

## Algorithm

1. Initialize weights randomly  $\sim \mathcal{N}(0, \sigma^2)$
2. Loop until convergence:
3. Compute gradient,  $\frac{\partial J(W)}{\partial W}$  *Not feasible to compute over all dataset*
4. Update weights,  $\mathbf{W} \leftarrow \mathbf{W} - \eta \frac{\partial J(W)}{\partial W}$
5. Return weights

# Gradient Descent

## Algorithm

1. Initialize weights randomly  $\sim \mathcal{N}(0, \sigma^2)$
2. Loop until convergence:
3. Compute gradient,  $\frac{\partial J(W)}{\partial W}$  *Compute over a mini-batch*
4. Update weights,  $\mathbf{W} \leftarrow \mathbf{W} - \eta \frac{\partial J(W)}{\partial W}$
5. Return weights

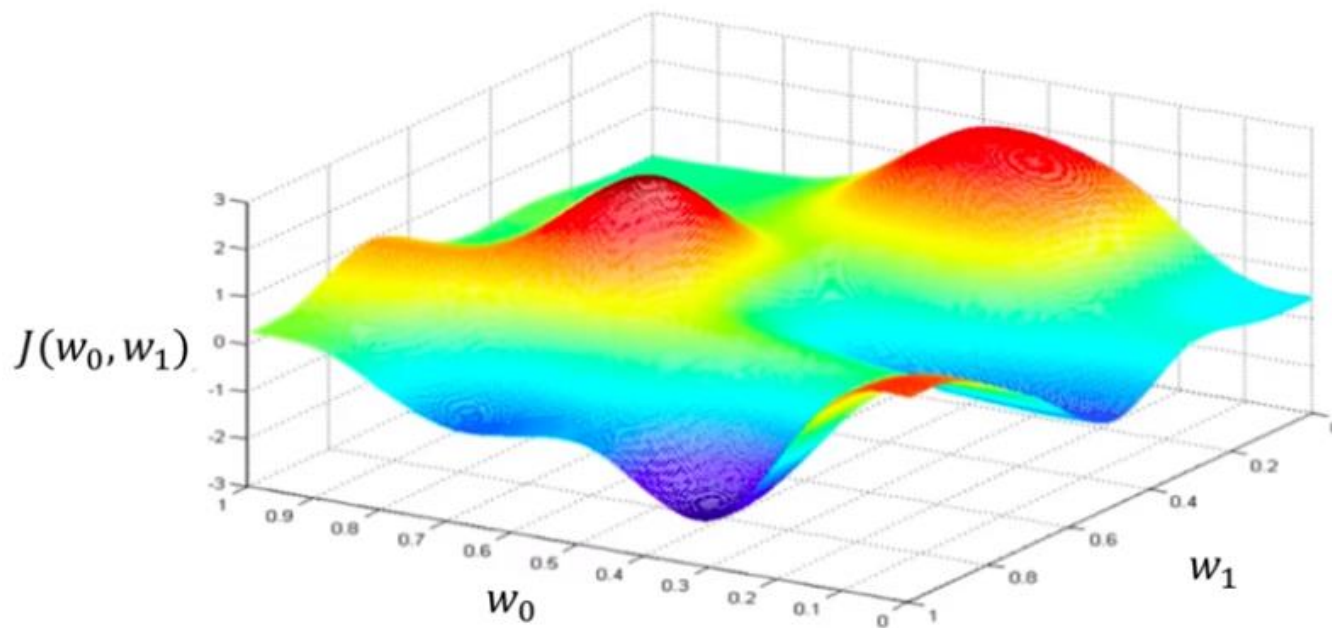
# Gradient Descent

## Algorithm

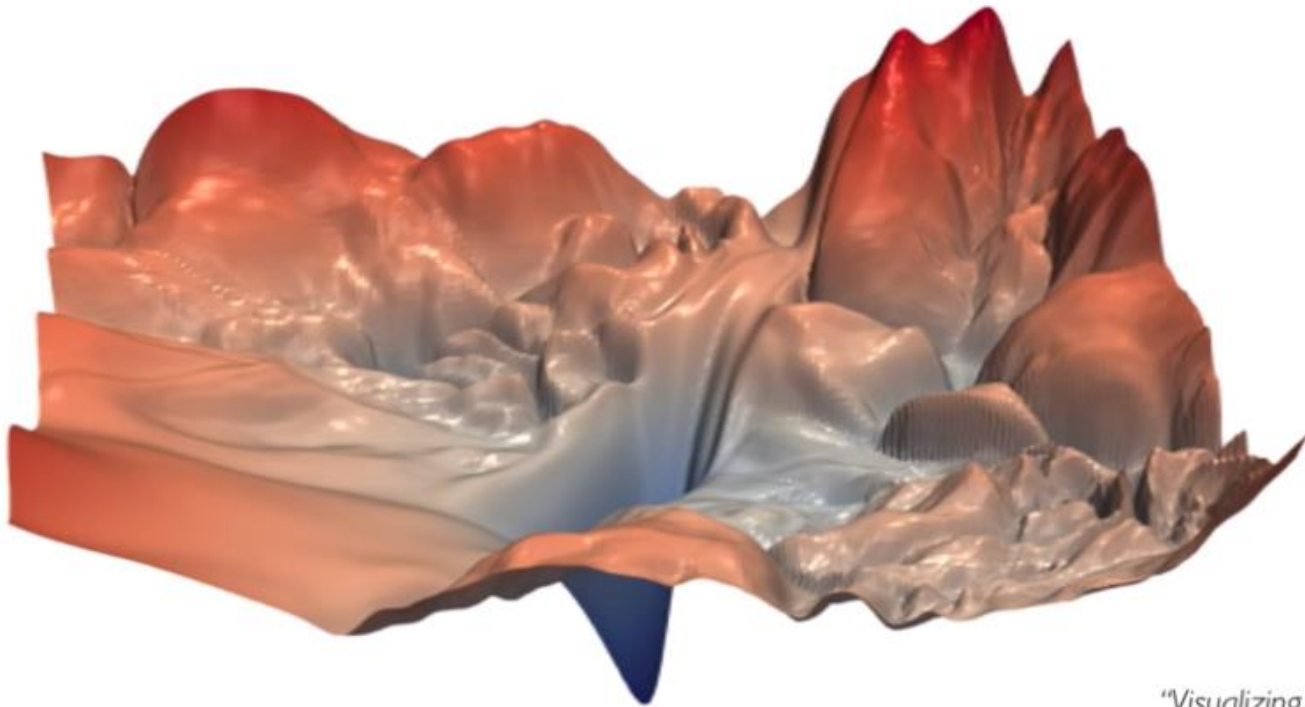
1. Initialize weights randomly  $\sim \mathcal{N}(0, \sigma^2)$
2. Loop until convergence:
3. Compute gradient,  $\frac{\partial J(W)}{\partial W}$  *Compute over a mini-batch*
4. Update weights,  $\mathbf{W} \leftarrow \mathbf{W} - \eta \frac{\partial J(W)}{\partial W}$
5. Return weights

*Parallelization: Batches can be split onto multiple GPUs*

# Loss/Cost Function



# Landscape Visualization

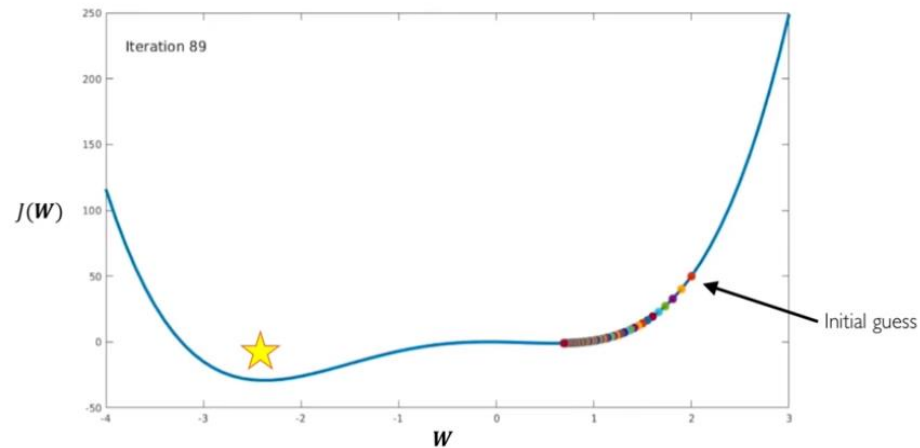


*"Visualizing the loss landscape  
of neural nets". Dec 2017.*



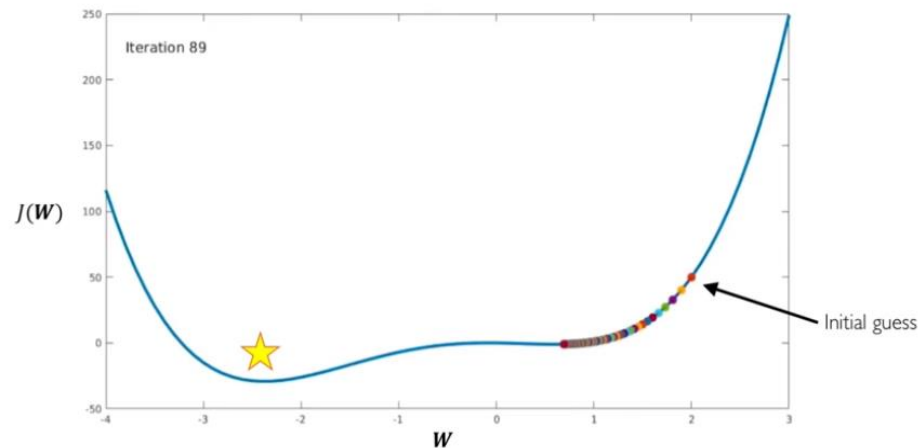
# Setting the Learning Rate

*Small learning rate converges slowly and gets stuck in false local minima*

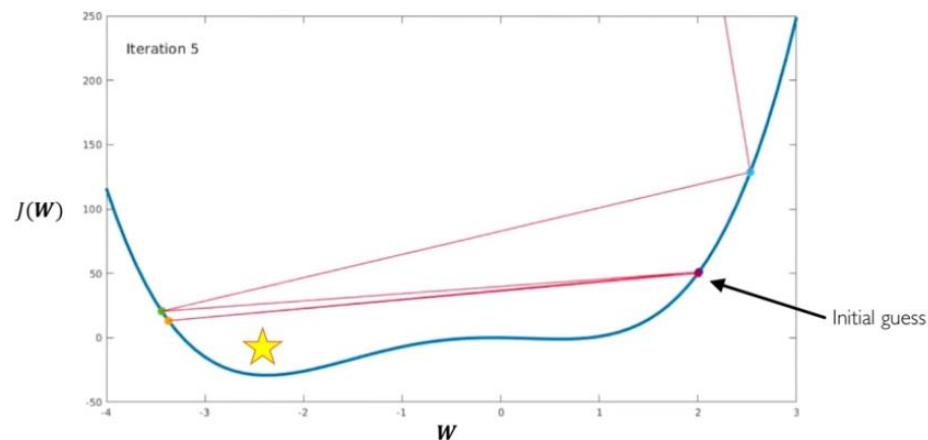


# Setting the Learning Rate

*Small learning rate converges slowly and gets stuck in false local minima*



*Large learning rates overshoot, become unstable and diverge*



# Setting the Learning Rate

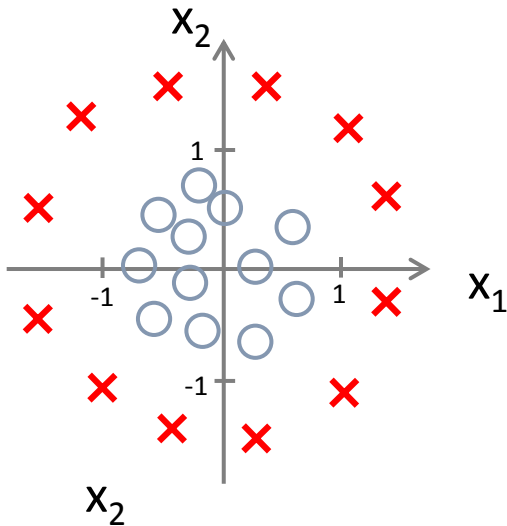
- How to select the learning Rate?
  - Try several, and see which works best
  - Start with a learning rate, and change it ***adaptively*** as the model trains
  - Many are implemented in Neural Network Tools

# Neural Networks Learn Features

logistic regression unit == **artificial neuron**

chain several units together == **neural network**

“earlier” units learn non-linear feature transformation

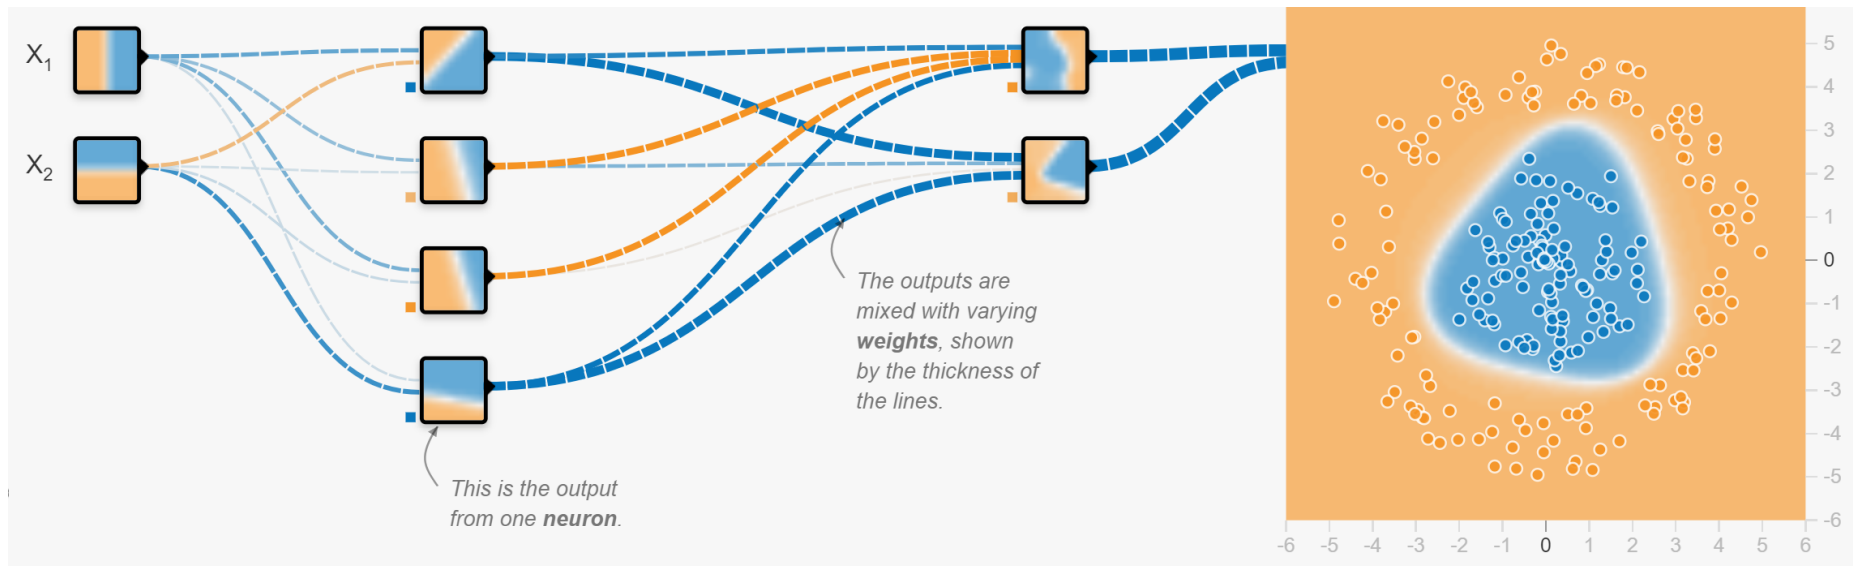


$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2)$$

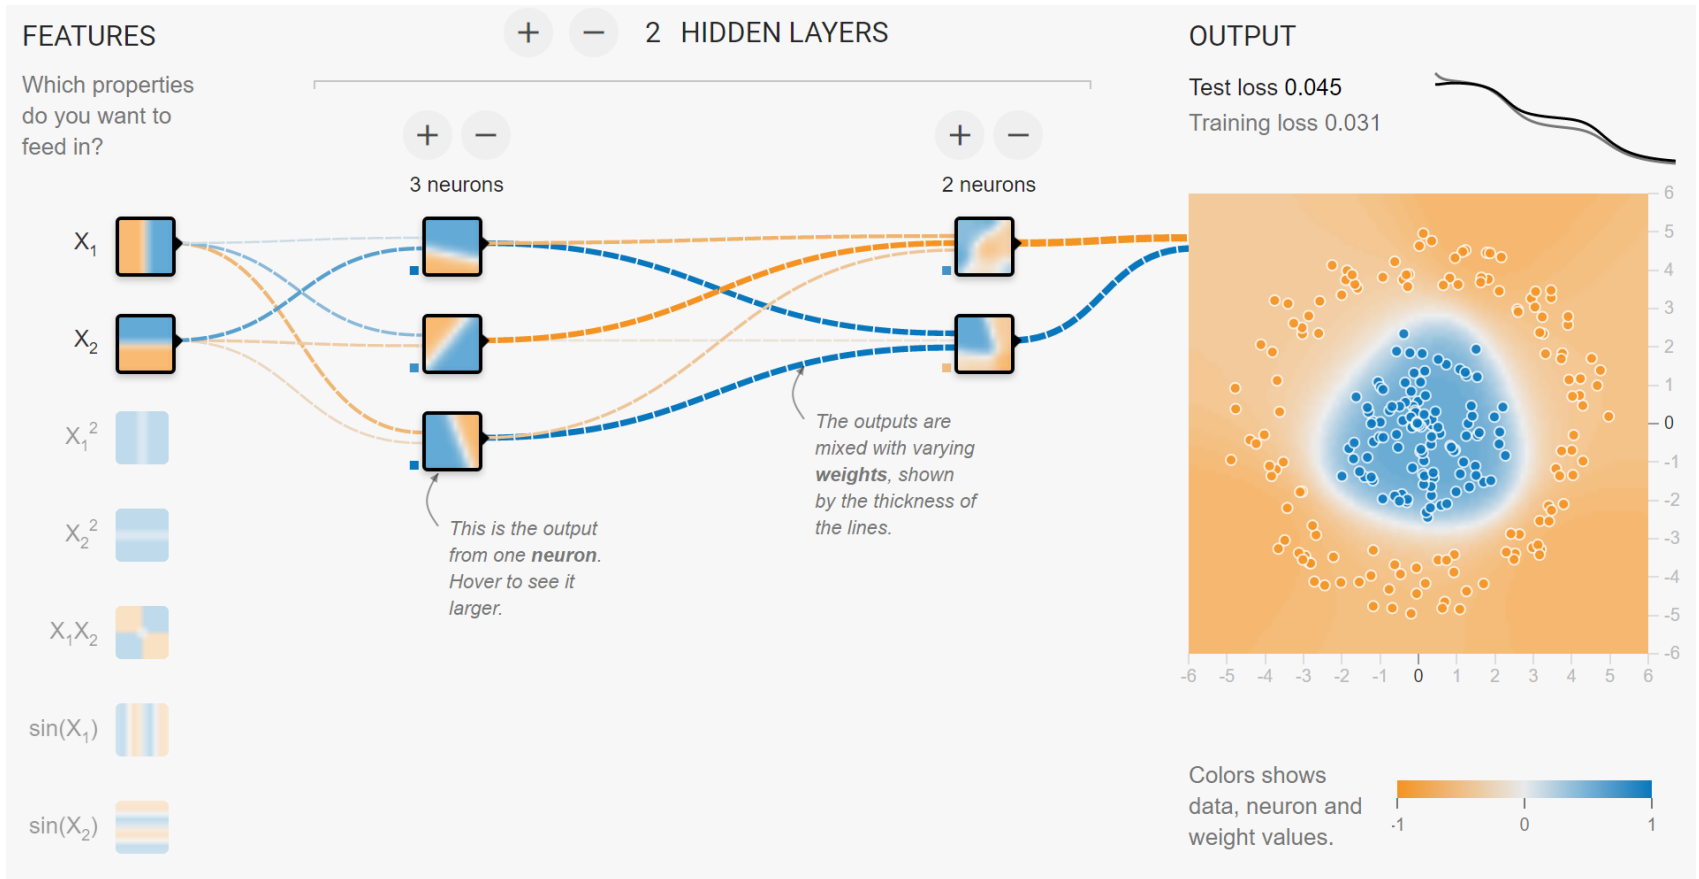
simple neural network

$$h(x) = g(\theta + \theta_1 h^{(1)}(x) + \theta_2 h^{(2)}(x) + \theta_3 h^{(3)}(x))$$

# Example



# Training a neural net: Demo



[Tensorflow playground](#)



# Deep Learning

## Architectures



# What is Deep Learning?

## ARTIFICIAL INTELLIGENCE

Any technique that enables computers to mimic human behavior



## MACHINE LEARNING

Ability to learn without explicitly being programmed



## DEEP LEARNING

Extract patterns from data using neural networks

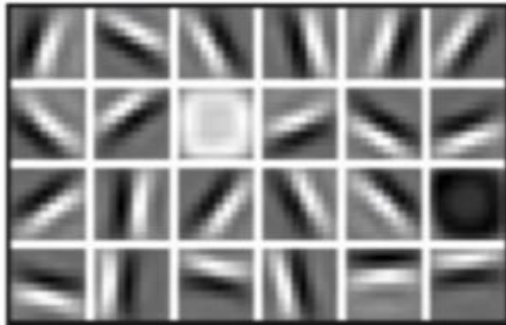
3 1 3 4 7 2  
1 7 4 2 3 5

# Why Deep Learning?

Hand engineered features are time consuming, brittle, and not scalable in practice

Can we learn the **underlying features** directly from data?

Low Level Features



Lines & Edges

Mid Level Features



Eyes & Nose & Ears

High Level Features



Facial Structure

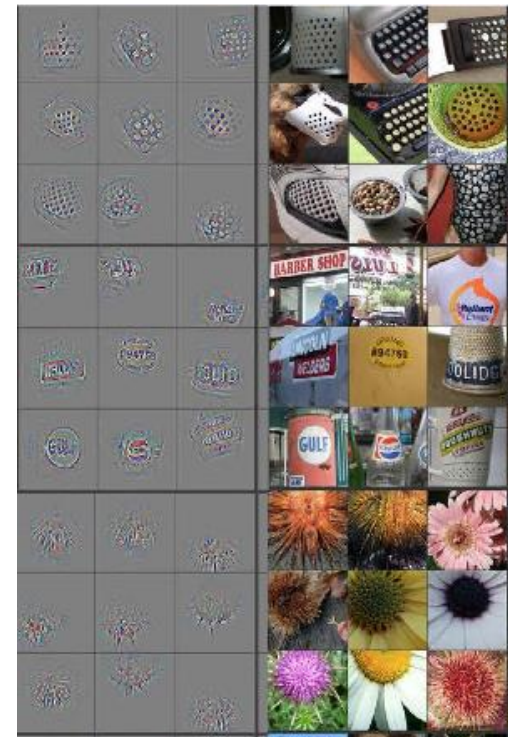
# Why Deep Learning?

## The Unreasonable Effectiveness of Deep Features



Maximal activations of pool<sub>5</sub> units

[R-CNN]

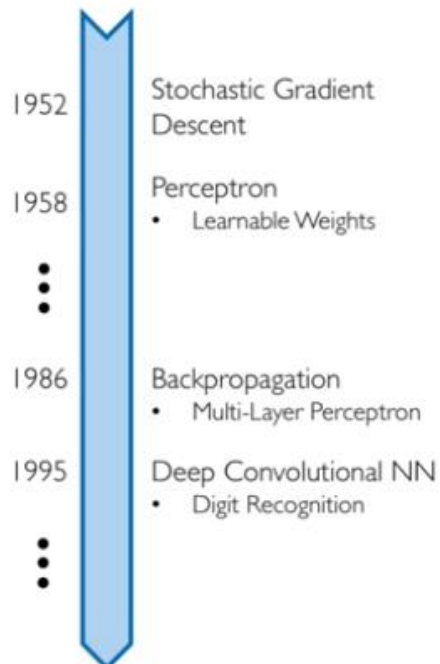


conv<sub>5</sub> DeConv visualization

[Zeiler-Fergus]

Rich visual structure of features deep in hierarchy.

# Why Now?



Neural Networks date back decades, so why the resurgence?

## 1. Big Data

- Larger Datasets
- Easier Collection & Storage

IMAGENET



## 2. Hardware

- Graphics Processing Units (GPUs)
- Massively Parallelizable



## 3. Software

- Improved Techniques
- New Models
- Toolboxes

