

Today: Outline

- **Neural networks cont'd:** architectures and learning; convolutional neural networks
- **Reminder:** Pre-lecture Material for Tue Sept 29




Neural Networks II

Architectures and Learning


Network architectures

Feed-forward


Fully connected



Recurrent



Convolutional






Neural Networks III

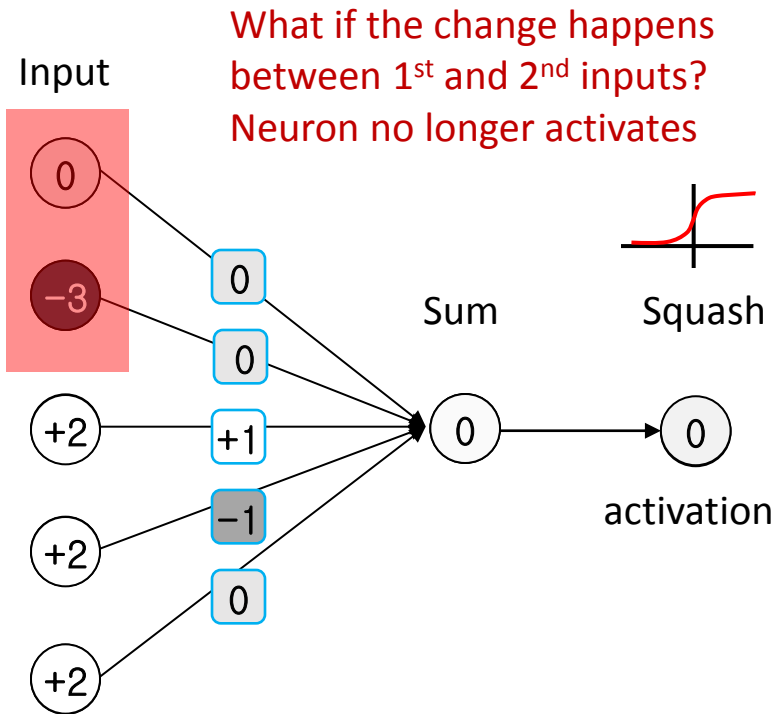
Convolutional Architectures

Multiplication vs convolution



- Recall, a neuron can be thought of as learning to spot certain features in the input
- E.g., this neuron detects change from high to low (light to dark) between 3rd and 4th inputs


Multiplication vs convolution



- Must have a new neuron for each new location of pattern???
- This is not efficient
- Solution: use convolution instead of multiplication

Multiplication vs convolution


Input



- New weights are of size 2×1 ; called **filter**, or **kernel**
- New output is the size of input minus 1 because of **boundary**
- New convolutional neurons all **share** the same weights! This is much more efficient; we learn the weights once instead of many times for each position

Multiplication vs convolution


Padded
Input



- New output is the size of input minus 1 because of **boundary**
- We can fix the boundary effect by **padding** the input with 0 and adding one more neuron

Multiplication vs convolution




Padded
Input



- Note, we move the filter by 1 each time, this is called **stride**

Multiplication vs convolution


Padded
Input



- Note, we move the filter by 1 each time, this is called **stride**
- Stride can be larger, e.g. here is stride 2

Multiplication vs convolution

Padded
Input

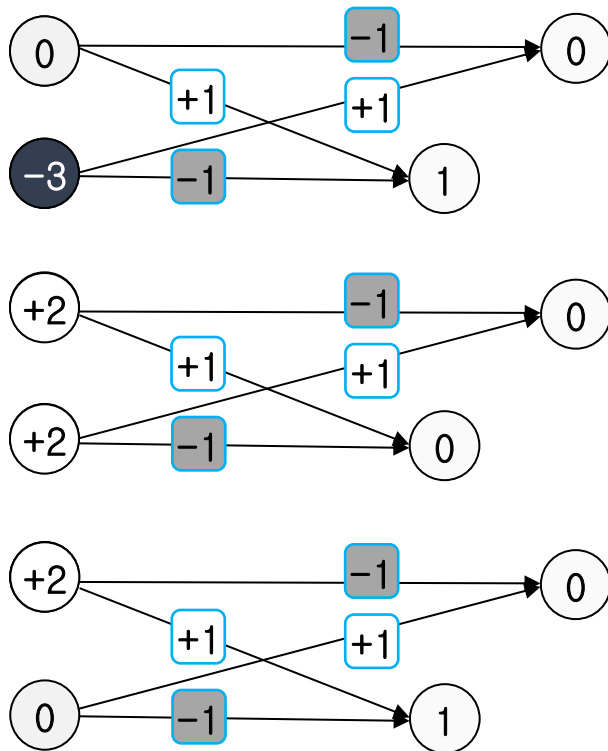


To summarize, this layer has

- Input 5×1 , padded to 6×1
- Kernel 2×1 with weights $[+1, -1]$
- Stride 2
- Output 3×1


Multiplication vs convolution

Padded
Input




- We can add another filter, this time to detect opposite change with weights $[-1 \ 1]$
- Unique filters are called **channels**

Multiplication vs convolution



- We can add another filter, this time to detect opposite change with weights $[-1 \ +1]$
- Unique filters are called **channels**

Multiplication vs convolution



- We can add another filter, this time to detect opposite change with weights $[-1 \ +1]$
- Unique filters are called **channels**

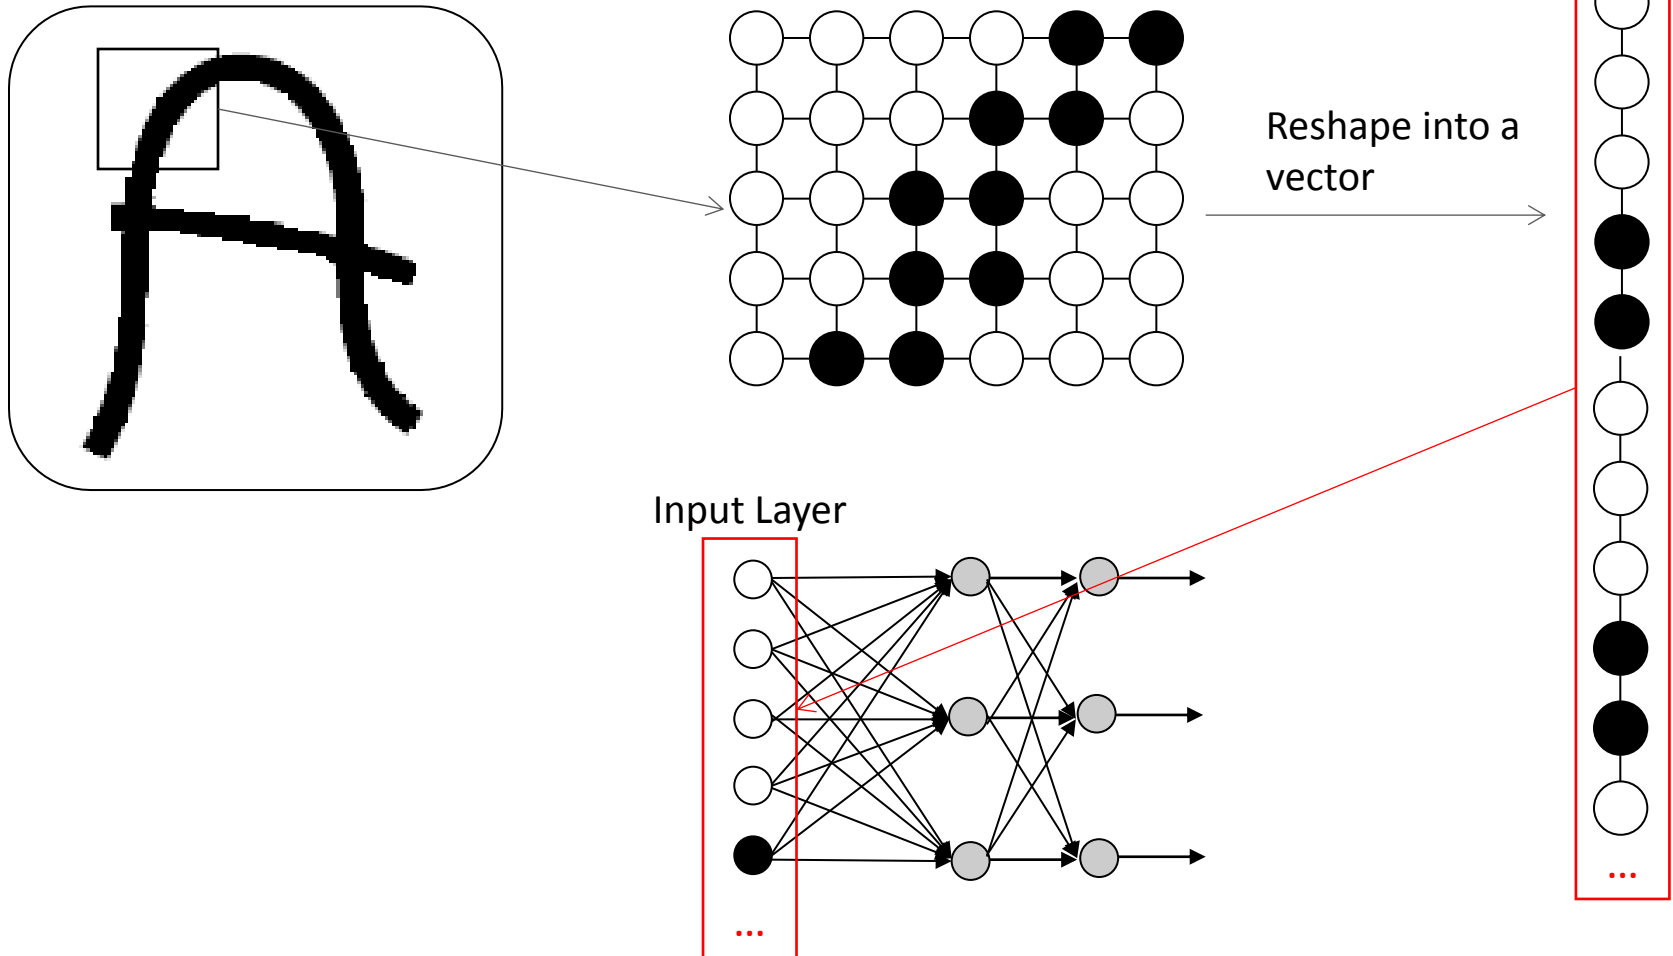


Convolutional Neural Networks

For images and other 2-D signals


Representing images

Fully connected




Convolution layer in 2D


- detect the same feature at different positions in the input, e.g. image
- preserve input topology



Convolution layer in 2D





Input image



Output map

Convolution layer in 2D



$$a = f(w_{11}x_{11} + w_{12}x_{12} + w_{13}x_{13} + \cdots w_{33}x_{33})$$

What weights correspond to these output maps?



*These are output maps
before thresholding*

*Hint: filters look like the
input they fire on*





$$\frac{\partial f}{\partial x}(x, y)$$

$$\partial x$$



$$\frac{\partial f}{\partial y}(x, y)$$

$$\partial y$$



Where is Waldo?



Input



filter

What will the output map look like?



Input



filter

What will the output map look like?



filter

Output

Here is Waldo




Input



filter


Stacking convolutional layers

- Each layer outputs multi-channel **feature maps** (like images)
- Next layer learns filters on previous layer's feature maps




Pooling layers

- Convolution with stride > 1 reduces the size of the input
- Another way to downsize the feature map is with **pooling**
- A pooling layer subsamples the input in each sub-window
 - **max-pooling**: chose the max in a window
 - **mean-pooling**: take the average




Pooling layer

- the **pooling** layers reduce the spatial resolution of each feature map
- Goal is to get a certain degree of shift and distortion **invariance**

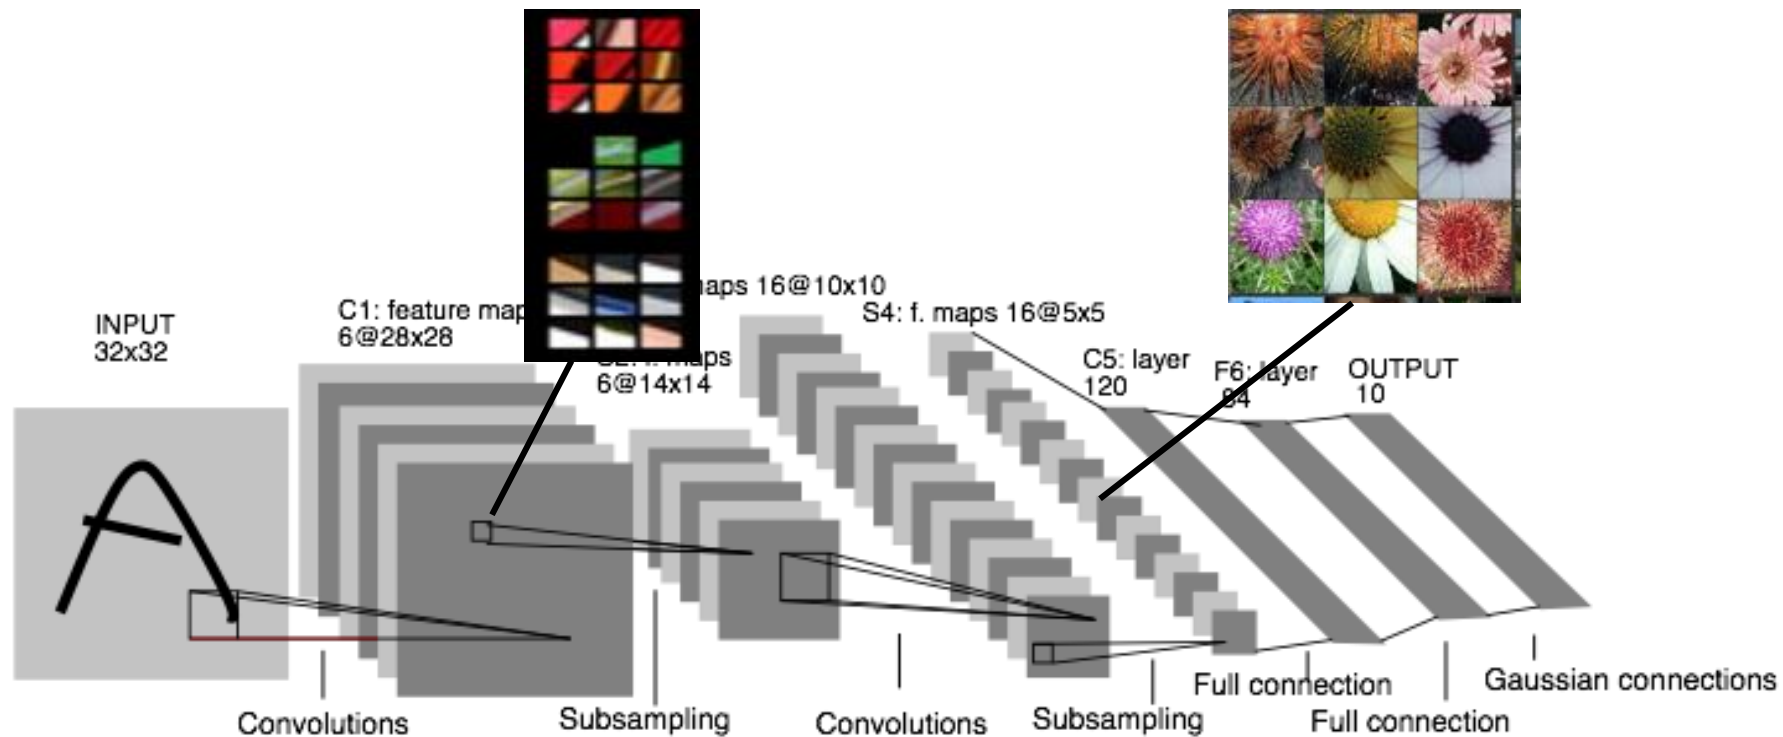


Putting it all together...



Convolutional Neural Network

A CNN is a better architecture for 2D signals



LeNet

Engineering

Features -> Architectures








Convolutional Neural Nets

Example


CIFAR-10 Demo ConvJS Network



ReLU: rectified linear unit



ReLU function $g(x) = \max(0, x)$




input (32x32x3)



filter size 5x5x3, stride 1





filter size 5x5x3, stride 1





input (32x32x3)



RELU

conv (32x32x16) params: $16 \times 5 \times 5 \times 3 + 16 = 1216$






input (32x32x3)




filter size 5x5x3, stride 1



conv (32x32x16) params: $16 \times 5 \times 5 \times 3 + 16 = 1216$





input (32x32x3)





pool (16x16x16)
pooling size 2x2, stride 2

filter size 5x5x3, stride 1



conv (32x32x16) params: $16 \times 5 \times 5 \times 3 + 16 = 1216$






input (32x32x3)





pool (16x16x16)
pooling size 2x2, stride 2

filter size 5x5x16, stride 1



RELU




pool (8x8x20)
pooling size 2x2, stride 2


filter size 5x5x3, stride 1



conv (32x32x16) params: $16 \times 5 \times 5 \times 3 + 16 = 1216$



conv (16x16x20) params: $20 \times 5 \times 5 \times 16 + 20 = 8020$



input (32x32x3)



One more conv+RELU+pool:

conv (8x8x20)

filter size 5x5x20, stride 1

relu (8x8x20)

pool (4x4x20)

pooling size 2x2, stride 2

parameters: $20 \times 5 \times 5 \times 20 + 20 = 10020$

fc (1x1x10); parameters: $10 \times 320 + 10 = 3210$



softmax (1x1x10)



Dog


cat

Car

⋮

Testing the network

- Show top three most likely classes



<http://cs.stanford.edu/people/karpathy/convnetjs/demo/cifar10.html>