

Today: Outline

- **Pre-lecture Material**
- **Recurrent Neural Networks**

- **Reminders:**

Problem Set 1, due: Oct 12 by midnight

Midterm Exam, in class, Oct 20 (Practice problems will be posted)

- **Announcement:**

No class on Oct 13 per BU Calendar
(Substitute Mon Schedule of Classes)



Neural Networks V

Pre-lecture Material


AI Generated Trailer

How did Watson generate a movie trailer?

- ☐ Using human expertise only
- ☐ Using tools humans use to generate trailers
- ☐ By identifying salient regions in the movie
- ☐ By detecting times where actions and emotions are predicted

AI Generated Trailer

- Analyze a movie and generate a trailer automatically
- How?
Detecting salient moments
e.g. action/emotions



Detecting Salient Regions

- Two sample actions:

Handstand Walking



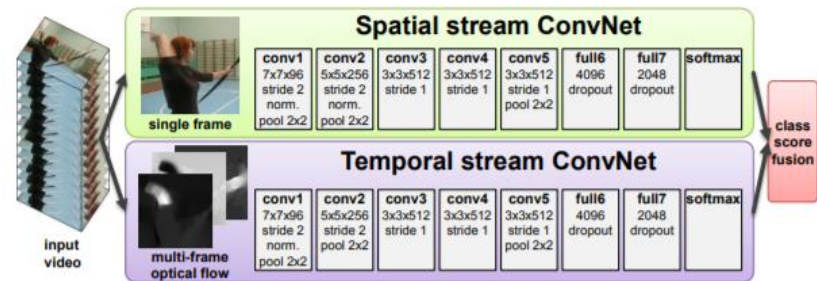
Ice Dancing



AI Generated Match Highlights

- IBM produces the official match highlights of Wimbledon and US Open tennis tournaments.
- https://www.usopen.org/en_US/video/2017-08-31/1504233424.html


- Multi-modal System
- Bias Considerations




Network architectures

Feed-forward


Fully connected



Convolutional



Recurrent





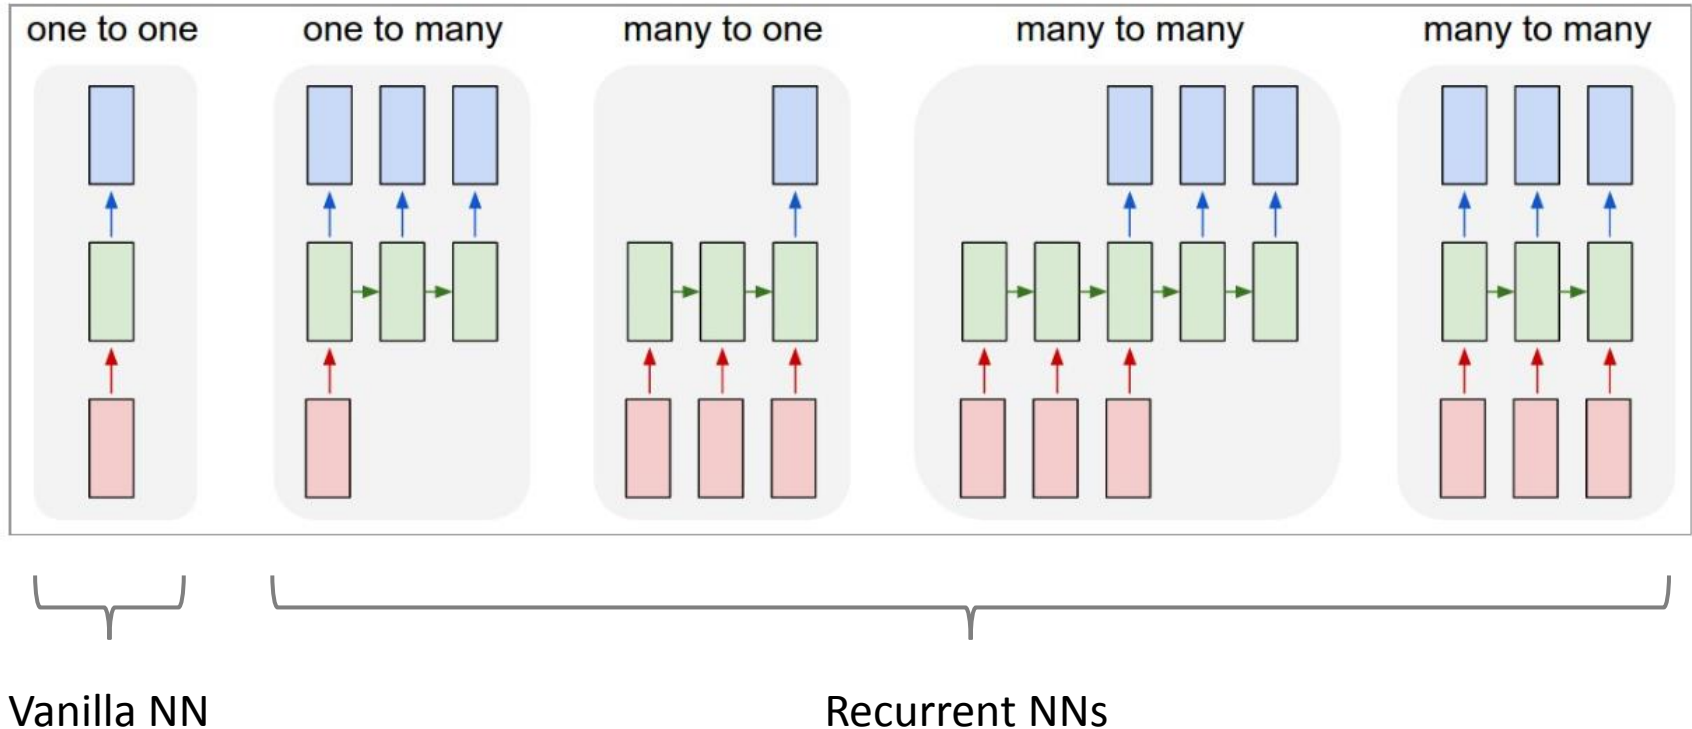
Neural Networks V

Recurrent Neural Networks


Sequential Data

- Sequences in our world:
 - Audio
 - Text
 - Video
 - Weather
 - Stock market
- RNNs are tools for making predictions about sequences.

Recurrent Neural Networks



One-to-one




Vanilla mode of processing without RNN

Example: Image classification

Example: One-to-one

Vanilla mode of processing without RNN

Example: Image classification




Prediction: "Dog"


Vanilla NN

Image

Recurrent Neural Networks



One-to-many



Sequence output


Example: Image captioning

Example: One-to-many


Sequence output

Example: Image Captioning

Bus driving down a snowy road next to trees <EOS>



Many-to-one




Sequence input

Examples: Sentiment analysis
Action recognition


Example: Many-to-one

Sequence input

Example: Sentiment analysis



Many-to-many



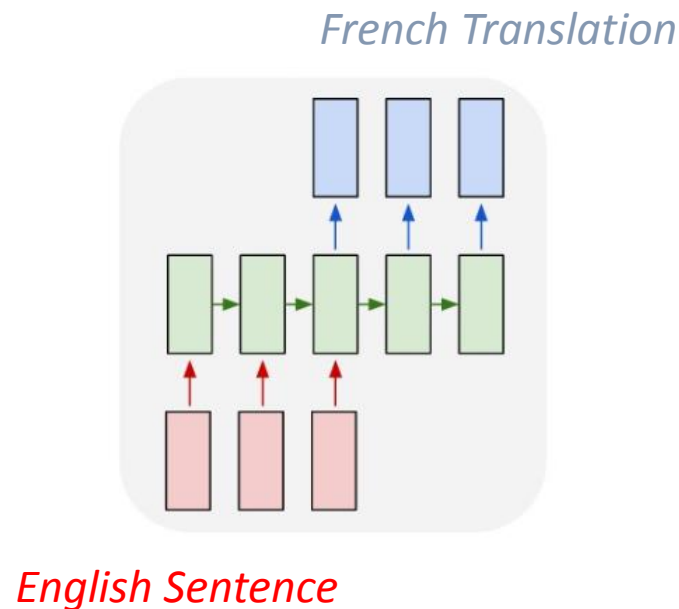
Sequence input and sequence output

Example: Machine translation


Example: Many-to-many

Sequence input and sequence output

Example: Machine translation



Synced Many-to-many




Synced sequence input and output

Examples: Tracking

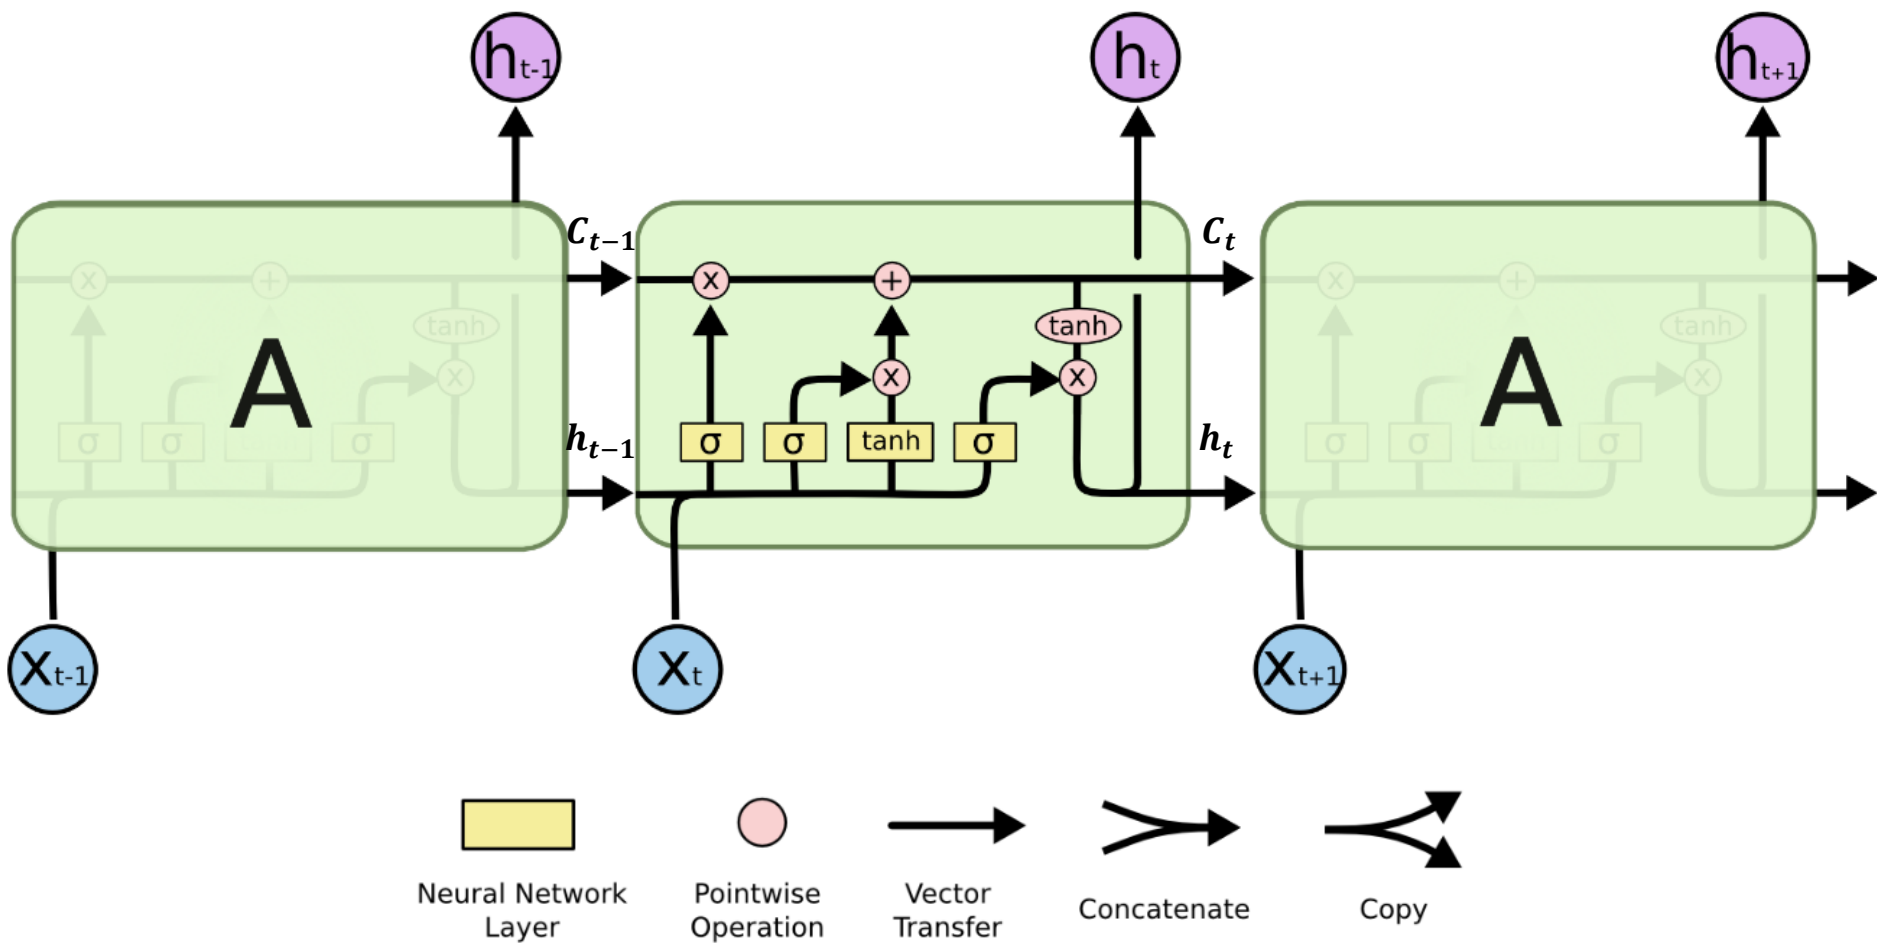
Early action detection

RNNs

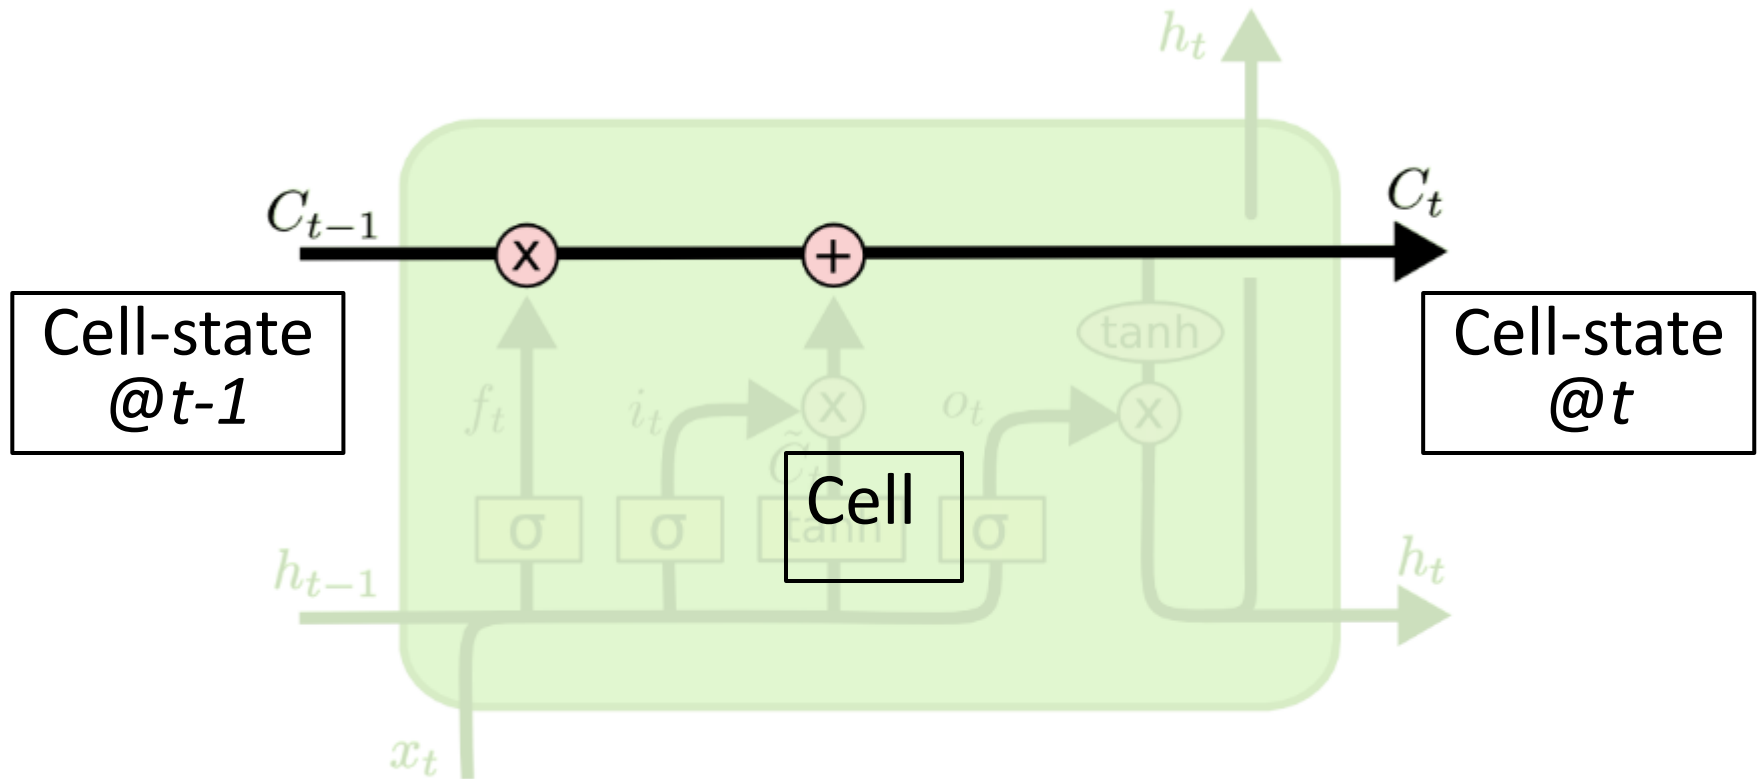
- In a standard RNN the repeating module has a simple structure. Example:



LSTMs




LSTM Memory / Cell State




Gate

- Composed of a sigmoid neural net layer and a pointwise multiplication operation.




Gate

- sigmoid: outputs numbers between:
 - zero “let nothing through,” and
 - one, “let everything through!”
- Example:



Gate

- sigmoid: outputs numbers between:
 - zero “let nothing through,” and
 - one, “let everything through!”
- Example:






Neural Networks VI

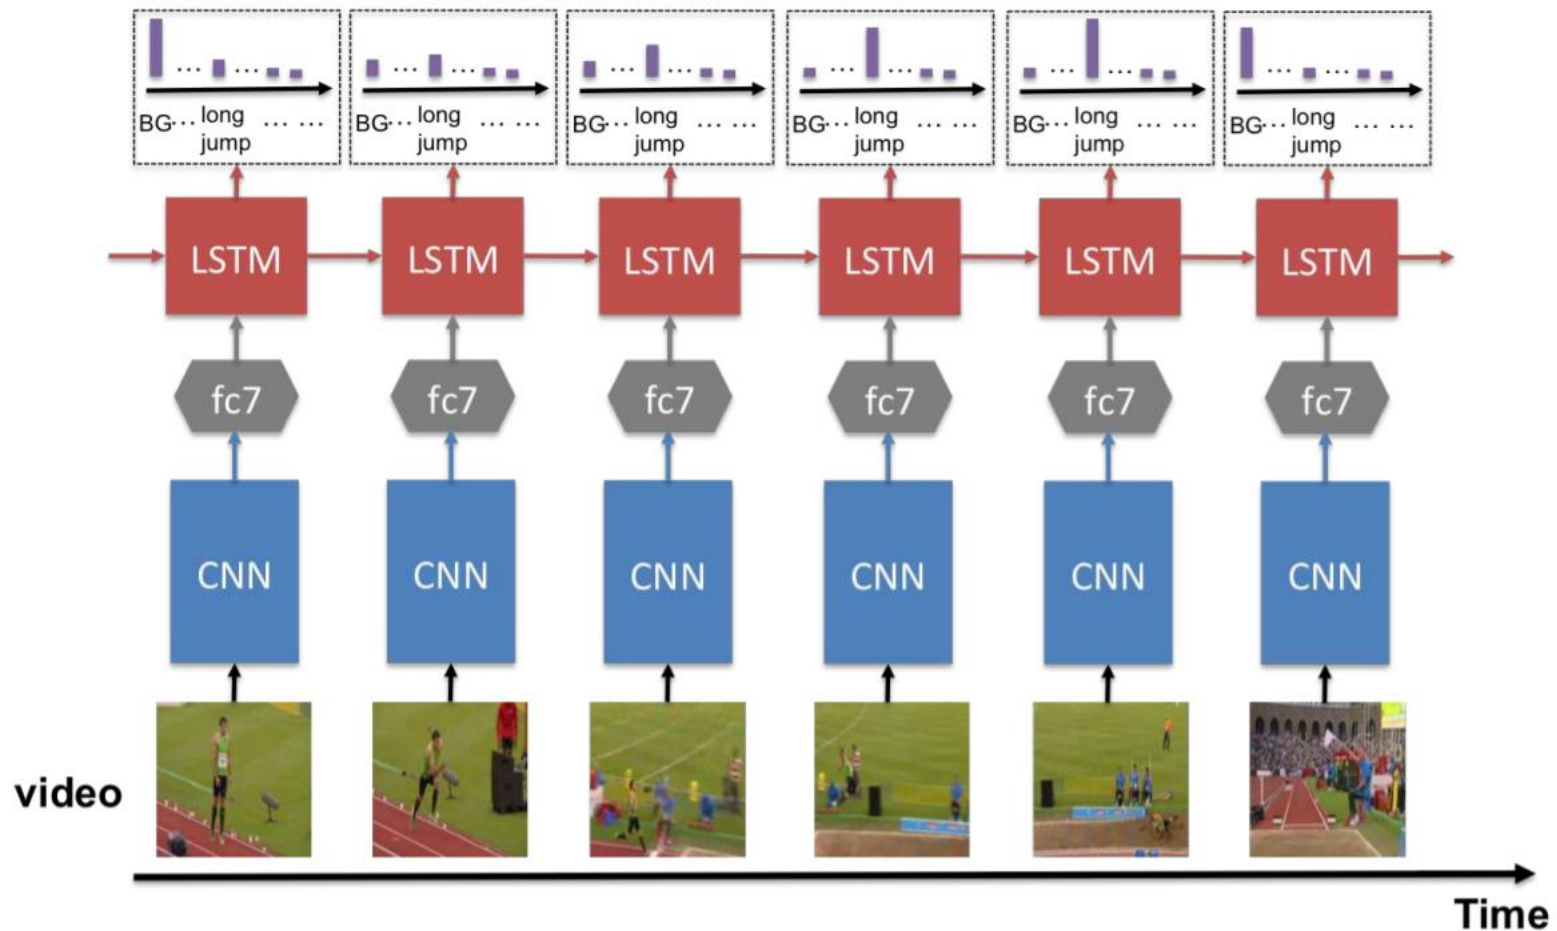
Applications of Recurrent Networks

Application 1: Video Classification

- CP: conditional class probability
- $\text{frame } i$ could be a feature describing frame i , example: CNN feature




Application 1: Video Classification



Application 2: Self-Driving Cars

- SA: steering angle
- $\text{frame } i$ could be a feature describing frame i ,
example: 3D-CNN feature



Application 2: Self-Driving Cars

- DeepTesla



Application 2: Self-Driving Cars

- Udacity winning team: *Team Komanda*
 - x_t : 3D convolution of image sequence
 - h_t : steering angle, speed, torque




Application 3: Character RNN

Character-level
language model
example

Vocabulary:
[h,e,l,o]

Example training
sequence:
“**hello**”




Application 3: Character RNN

Character-level language model example

Vocabulary:
[h,e,l,o]

Example training
sequence:
“hello”

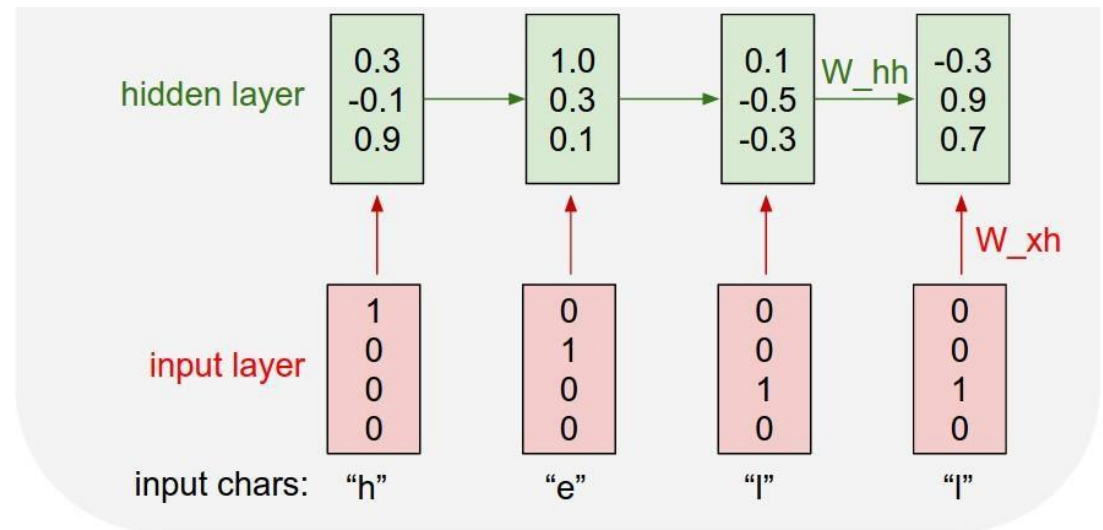


Application 3: Character RNN

Character-level language model example

Vocabulary:
[h,e,l,o]

Example training
sequence:
“hello”

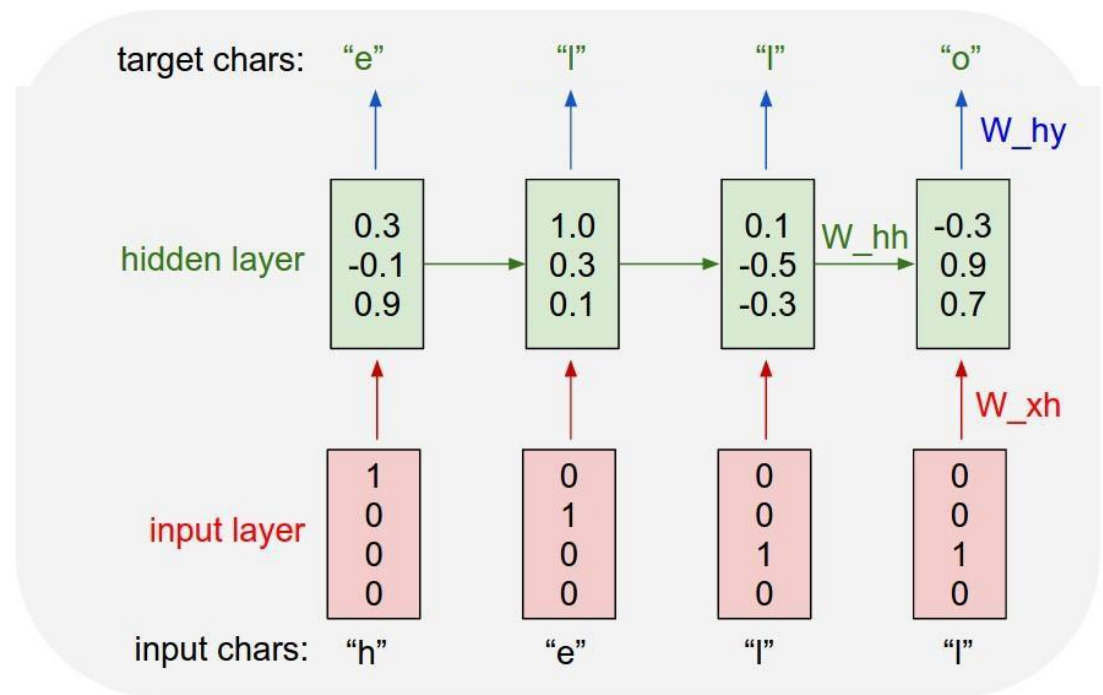


Application 3: Character RNN

Character-level language model example

Vocabulary:
[h,e,l,o]

Example training
sequence:
“hello”



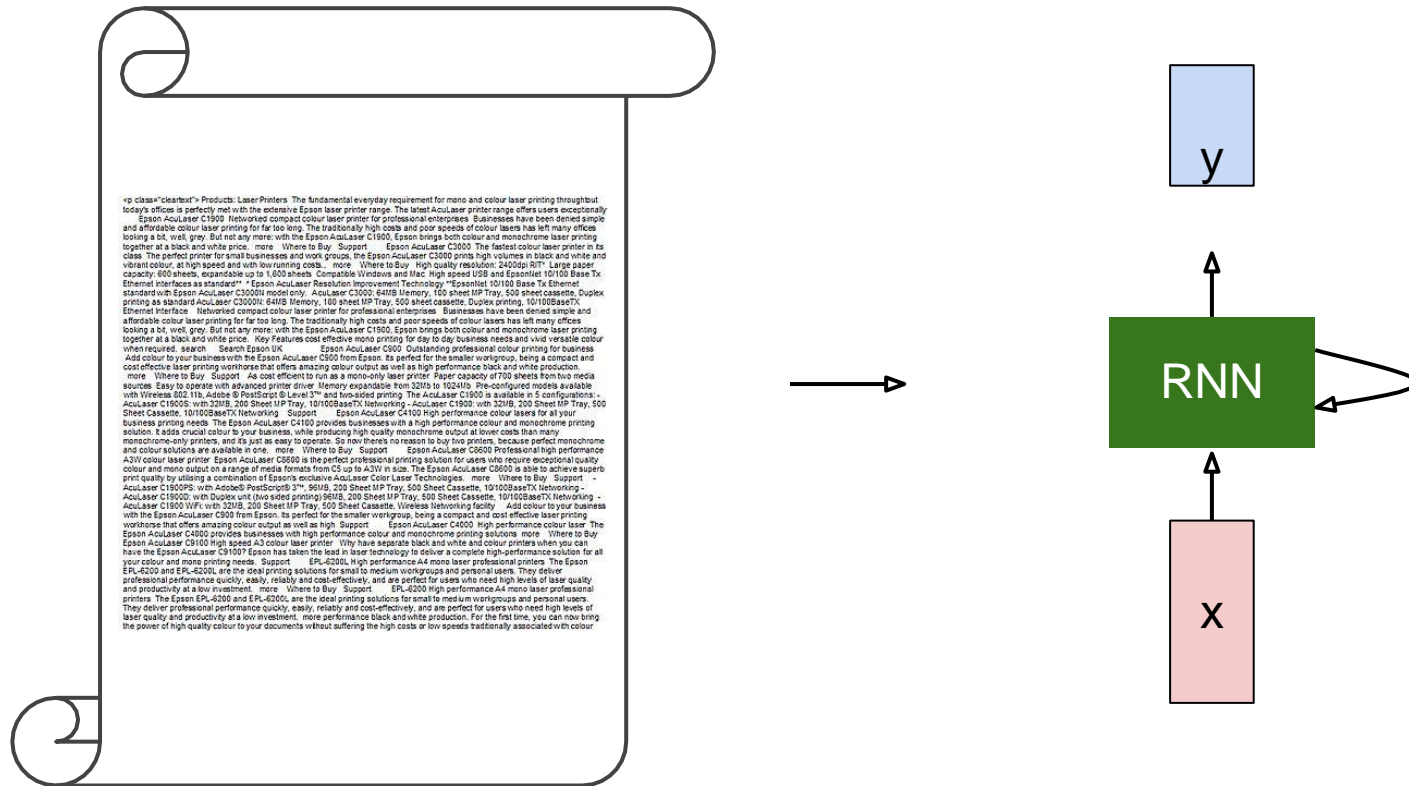
Application 4: Reading cursive

handwriting

- This is a natural task for an RNN.
- The input is a sequence of (x,y,p) coordinates of the tip of the pen, where p indicates whether the pen is up or down.
- The output is a sequence of characters.
- Graves & Schmidhuber (2009) showed that RNNs with LSTM are currently the best systems for reading cursive writing.
 - They used a sequence of small images as input rather than pen coordinates.

Application 5: StyleText Generation

Training text: William Shakespeare



Fei-Fei Li & Andrej Karpathy & Justin Johnson

Application 5: StyleText Generation

at
first:

tyntd-iafhatawiaoihrdemot lytdws e ,tfti, astai f ogoh eoase rrranbyne 'nhthnee e
plia tkllrgd t o idoe ns,smtt h ne etie h,hregtrs niglike,aoaenns lng



train more

"Tmont thithey" fomesscerliund
Keushey. Thom here
sheulke, anmerenith ol sivh I lalterthend Bleipile shuw y fil on aseterlome
coaniogennc Phe lism thond hon at. MeiDimorotion in ther thize."



train more

Aftair fall unsuch that the hall for Prince Velzonski's that me of
her hearly, and behs to so arwage fiving were to it beloge, pavu say falling misfort
how, and Gogition is so overelical and ofter.



train more

"Why do what that day," replied Natasha, and wishing to himself the fact the
princess, Princess Mary was easier, fed in had oftended him.
Pierre aking his soul came to the packs and drove up his father-in-law women.

Application 6: Code Generation

Train on C code

The screenshot shows the GitHub interface for the 'torvalds / linux' repository. At the top, there's a search bar and navigation links like 'Explore', 'Gist', 'Blog', and 'Help'. The repository name 'torvalds / linux' is prominently displayed, along with statistics: 3,711 watches, 23,054 stars, and 9,141 forks. Below this, the 'Linux kernel source tree' is highlighted. A summary bar indicates 520,037 commits, 1 branch, 420 releases, and 5,039 contributors. The main content area shows a list of recent commits, with the most recent one by 'torvalds' 9 hours ago, titled 'Merge branch 'drm-fixes' of git://people.freedesktop.org/~airlied/linux'. The commit details show a merge of 'git://git.kernel.org/pub/scm/linux/kernel/git/nab/target-pending'. A table of recent commits follows, listing various merges and updates across different kernel components like 'Documentation', 'arch', 'block', 'crypto', 'drivers', 'firmware', 'fs', 'include', and 'init'. On the right side, there are links for 'Code', 'Pull requests' (74), 'Pulse', and 'Graphs'. At the bottom right, the 'HTTPS clone URL' is provided as 'https://github.com/torvalds/linux.git', along with options to 'Clone in Desktop' or 'Download ZIP'.

torvalds / linux

Linux kernel source tree

520,037 commits 1 branch 420 releases 5,039 contributors

branch: master - linux / +

Merge branch 'drm-fixes' of git://people.freedesktop.org/~airlied/linux

torvalds authored 9 hours ago latest commit 4b1786927d

Directory	Commit Message	Time Ago
Documentation	Merge git://git.kernel.org/pub/scm/linux/kernel/git/nab/target-pending	6 days ago
arch	Merge branch 'x86-urgent-for-linus' of git://git.kernel.org/pub/scm/l...	a day ago
block	block: discard bdi_unregister() in favour of bdi_destroy()	9 days ago
crypto	Merge git://git.kernel.org/pub/scm/linux/kernel/git/herbert/crypto-2.6	10 days ago
drivers	Merge branch 'drm-fixes' of git://people.freedesktop.org/~airlied/linux	9 hours ago
firmware	firmware/hex2fw.c: restore missing default in switch statement	2 months ago
fs	vfs: read file_handle only once in handle_to_path	4 days ago
include	Merge branch 'perf-urgent-for-linus' of git://git.kernel.org/pub/scm/...	a day ago
init	init: fix regression by supporting devices with major:minor:offset fo...	a month ago

HTTPS clone URL
https://github.com/torvalds/linux.git

You can clone with HTTPS, SSH, or Subversion.

Clone in Desktop

Download ZIP

Application 6: Code Generation

```
static void do_command(struct seq_file *m, void *v)
{
    int column = 32 << (cmd[2] & 0x80);
    if (state)
        cmd = (int)(int_state ^ (in_8(&ch->ch_flags) & Cmd) ? 2 : 1);
    else
        seq = 1;
    for (i = 0; i < 16; i++) {
        if (k & (1 << 1))
            pipe = (in_use & UMXTHREAD_UNCCA) +
                ((count & 0x00000000ffffffff8) & 0x0000000f) << 8;
        if (count == 0)
            sub(pid, ppc_md.kexec_handle, 0x20000000);
        pipe_set_bytes(i, 0);
    }
    /* Free our user pages pointer to place camera if all dash */
    subsystem_info = &of_changes[PAGE_SIZE];
    rek_controls(offset, idx, &soffset);
    /* Now we want to deliberately put it to device */
    control_check_polarity(&context, val, 0);
    for (i = 0; i < COUNTER; i++)
        seq_puts(s, "policy ");
}
```

Generated
C code

Application 7: Writing a Movie Script



<https://arstechnica.com/the-multiverse/2016/06/an-ai-wrote-this-movie-and-its-strangely-moving/>